

# AMD Sinkclose

## Universal SMM Privilege Escalation

Enrique Nissim      @kiqueNissim


Krzysztof Okupski      @exminium





# Outline

- Technical background
  - Privilege levels and SMM security
  - Remapping attacks
- Exploitation
  - Exploit development
  - Demo
- Attack paths
- Conclusions





# SMM Introduction




# Introducing System Management Mode

- One of the most powerful execution modes in x86
  - Full access to system and I/O device memory
  - Access to the SPI flash (potential for persistence)
- Invisible to the rest of the system
  - Hidden from the OS and Hypervisor
  - EDRs cannot help here



# Privilege levels






Ring 3

Ring 0

Ring -2




Boot-time

Run-time

# System Management Interrupts

- SMM is entered using a special external interrupt called the system-management interrupt (SMI)
- After an SMI is received by the processor, the processor saves the processor state in a separate address space, called System Management RAM (SMRAM)








# Previous research

- Blogs
  - [Exploring the security configuration of AMD platforms](#) (2022)
  - [Adventures in the Platform Security Coordinated Disclosure Circus](#) (2023)
  - [Back to the Future with Platform Security](#) (2023)
  - [Exploring AMD Platform Secure Boot](#) (2023)
- Couple of CVEs

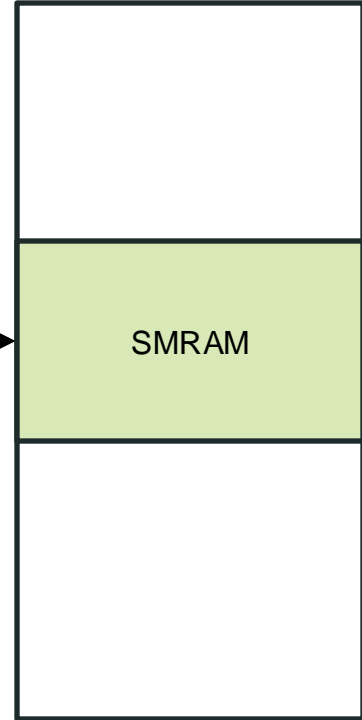

<i>CVE-2023-20576</i>	<i>CVE-2023-20577</i>	<i>CVE-2023-20579</i>
<i>CVE-2023-20587</i>	<i>CVE-2023-20596</i>	<i>CVE-2023-31100</i>
<i>CVE-2023-28468</i>	<i>CVE-2023-2290</i>	<i>CVE-2023-5078</i>
- Tooling: <https://github.com/IOActive/Platbox>



# SMM Security





CPU  
SMM




SMRAM

DRAM



- Execution disallowed
- Reads FFs
- Writes are discarded




DRAM




# TSEG Region


- How does the memory controller protect SMRAM?
  - At boot-time BIOS configures two registers to setup the TSEG Region

MSRC001\_0112 SMM TSeg Base Address (SMMAddr)



MSRC001\_0113 SMM TSeg Mask (SMMMask)







# Summary of SMRAM Registers

- MSRC001\_0111 (SMM\_BASE used for SMM base address)
- MSRC001\_0112 (SMM TSeg Base Address (SMMAddr))
- MSRC001\_0113 (SMM TSeg Mask (SMMMMask))
- MSRC001\_0015[SmmLock] (HWCR used for locking the config)

*These need to be configured for each core*



# Differences between AMD and Intel MSRs

- On Intel systems there are specific MSRs that are only accessible while the processor is executing at SMM
  - *Example: IA32\_SMBASE (SMM base register)*
  - *Obtaining this value could be considered a leak*
- On AMD all the MSRs that are related to the security of SMM are accessible from ring 0
  - *Note that when SmmLock bit is set, accesibility does not imply the configuration can be changed even from SMM*





# Spotting the bug

Bits	Description
63:40	Reserved.
39:17	<b>TSegMask[39:17]: TSeg address range mask.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only ELSE Read-write ENDIF. See <a href="#">MSRC001_0112</a> .
16:15	Reserved.
14:12	<b>TMTypDram: TSeg address range memory type.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the memory type for SMM accesses to the TSeg range that are directed to DRAM. See: <a href="#">Table 219 [Valid Values for Memory Type Definition]</a> .
11	Reserved.
10:8	<b>AMTypDram: ASeg Range Memory Type.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the memory type for SMM accesses to the ASeg range that are directed to DRAM. See: <a href="#">Table 219 [Valid Values for Memory Type Definition]</a> .
7:6	Reserved.
5	<b>TMTypIoWc: non-SMM TSeg address range memory type.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the attribute of TSeg accesses that are directed to MMIO space. 0=UC (uncacheable). 1=WC (write combining).
4	<b>AMTypIoWc: non-SMM ASeg address range memory type.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the attribute of ASeg accesses that are directed to MMIO space. 0=UC (uncacheable). 1=WC (write combining).
3	<b>TClose: send TSeg address range data accesses to MMIO.</b> Read-write. 1=When in SMM, direct data accesses in the TSeg address range to MMIO space. See <a href="#">AClose</a> .
2	<b>AClose: send ASeg address range data accesses to MMIO.</b> Read-write. 1=When in SMM, direct data accesses in the ASeg address range to MMIO space.  [A, T]Close allows the SMI handler to access the MMIO space located in the same address region as the [A, T]Seg. When the SMI handler is finished accessing the MMIO space, it must clear the bit. Failure to do so before resuming from SMM causes the CPU to erroneously read the save state from MMIO space.
1	<b>TValid: enable TSeg SMM address range.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only. ELSE Read-write. ENDIF. 1=The TSeg address range SMM enabled.
0	<b>AValid: enable ASeg SMM address range.</b> IF <a href="#">MSRC001_0015</a> [SmmLock] THEN Read-only. ELSE Read-write. ENDIF. 1=The ASeg address range SMM enabled.

Bits	Description
63:40	Reserved.
39:17	<b>TSegMask[39:17]: TSeg address range mask.</b> IF MSRC001_0015[SmmLock] THEN Read-only ELSE Read-write ENDIF. See MSRC001_0112.
16:15	Reserved.
14:12	<b>TMTypDram: TSeg address range memory type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the memory type for SMM accesses to the TSeg range that are directed to DRAM. See: Table 219 [Valid Values for Memory Type Definition].
11	Reserved.
10:8	<b>AMTypDram: ASeg Range Memory Type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the memory type for SMM accesses to the ASeg range that are directed to DRAM. See: Table 219 [Valid Values for Memory Type Definition].
7:6	Reserved.
5	<b>TMTypIoWc: non-SMM TSeg address range memory type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the attribute of TSeg accesses that are directed to MMIO space. 0=UC (uncacheable). 1=WC (write combining).
4	<b>AMTypIoWc: non-SMM ASeg address range memory type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the attribute of ASeg accesses that are directed to MMIO space. 0=UC (uncacheable). 1=WC (write combining).
3	<b>TClose: send TSeg address range data accesses to MMIO.</b> Read-write. 1=When in SMM, direct data accesses in the TSeg address range to MMIO space. See AClose.
2	<b>AClose: send ASeg address range data accesses to MMIO.</b> Read-write. 1=When in SMM, direct data accesses in the ASeg address range to MMIO space.  [A, T]Close allows the SMI handler to access the MMIO space located in the same address region as the [A, T]Seg. When the SMI handler is finished accessing the MMIO space, it must clear the bit. Failure to do so before resuming from SMM causes the CPU to erroneously read the save state from MMIO space.
1	<b>TValid: enable TSeg SMM address range.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. 1=The TSeg address range SMM enabled.
0	<b>AValid: enable ASeg SMM address range.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. 1=The ASeg address range SMM enabled.

Bits	Description
63:40	Reserved.
39:17	<b>TSegMask[39:17]: TSeg address range mask.</b> IF MSRC001_0015[SmmLock] THEN Read-only ELSE Read-write ENDIF. See MSRC001_0112.
16:15	Reserved.
14:12	<b>TMTypDram: TSeg address range memory type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the memory type for SMM accesses to the TSeg range that are directed to DRAM. See: Table 219 [Valid Values for Memory Type Definition].
11	Reserved.
10:8	<b>AMTypDram: ASeg Range Memory Type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the memory type for SMM accesses to the ASeg range that are directed to DRAM. See: Table 219 [Valid Values for Memory Type Definition].
7:6	Reserved.
5	<b>TMTypIoWc: non-SMM TSeg address range memory type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the attribute of TSeg accesses that are directed to MMIO space. 0=UC (uncacheable). 1=WC (write combining).
4	<b>AMTypIoWc: non-SMM ASeg address range memory type.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. Specifies the attribute of ASeg accesses that are directed to MMIO space. 0=UC (uncacheable). 1=WC (write combining).
3	<b>TClose: send TSeg address range data accesses to MMIO.</b> Read-write. 1=When in SMM, direct data accesses in the TSeg address range to MMIO space. See AClose.
2	<b>AClose: send ASeg address range data accesses to MMIO.</b> Read-write. 1=When in SMM, direct data accesses in the ASeg address range to MMIO space.  [A, T]Close allows the SMI handler to access the MMIO space located in the same address region as the [A, T]Seg. When the SMI handler is finished accessing the MMIO space, it must clear the bit. Failure to do so before resuming from SMM causes the CPU to erroneously read the save state from MMIO space.
1	<b>TValid: enable TSeg SMM address range.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. 1=The TSeg address range SMM enabled.
0	<b>AValid: enable ASeg SMM address range.</b> IF MSRC001_0015[SmmLock] THEN Read-only. ELSE Read-write. ENDIF. 1=The ASeg address range SMM enabled.





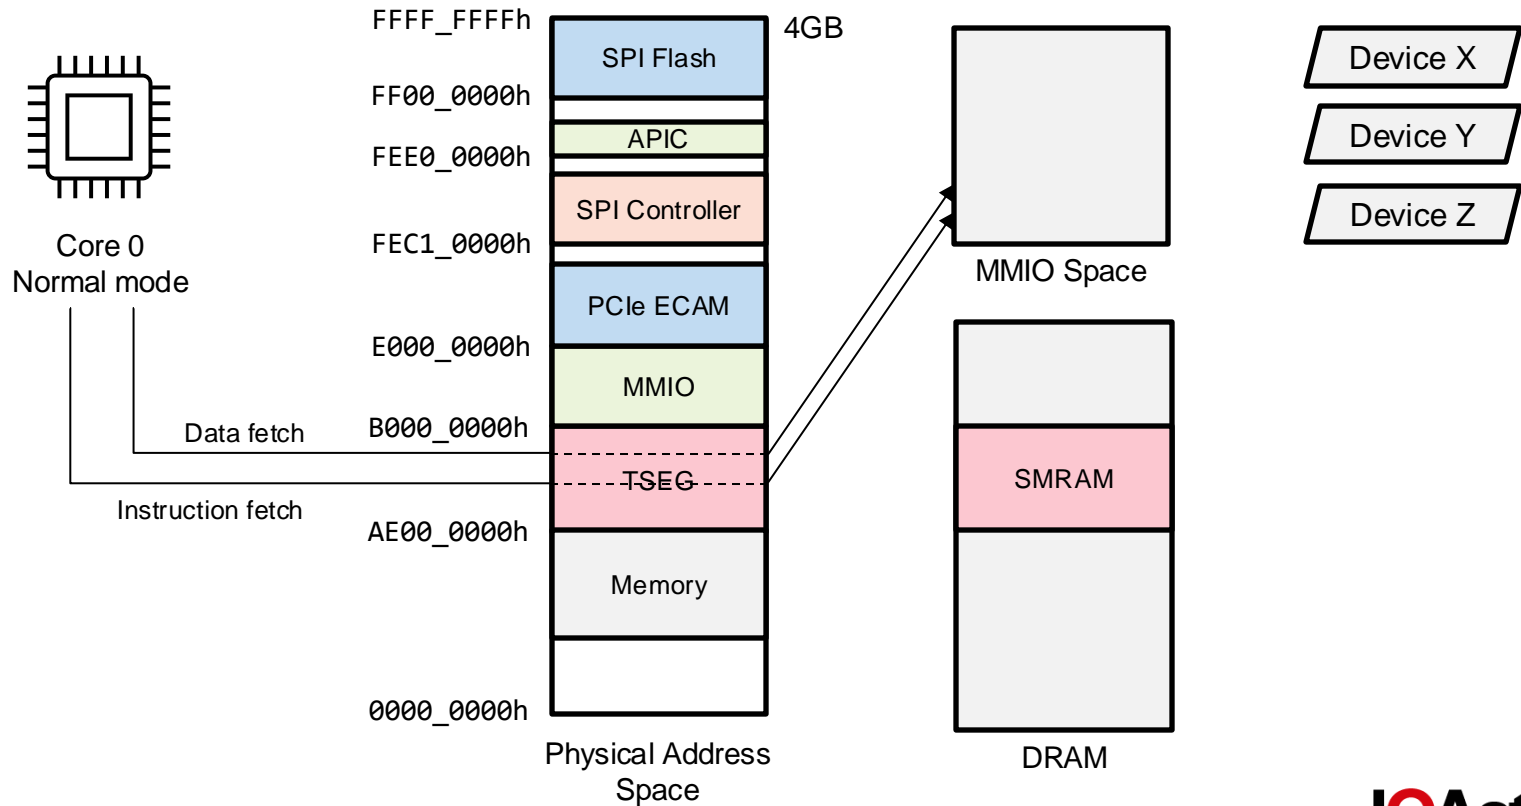
# MSR C001\_0113 SMM TSeg Mask (SMMMask)

This register specifies how accesses to the ASeg and TSeg address ranges are controlled as follows:

- If [A,T]Valid=1, then:
  - If in SMM, then:
    - If [A, T]Close=0, then the accesses are directed to DRAM with memory type as specified in [A, T]MTypeDram.
    - If [A, T]Close=1, then instruction accesses are directed to DRAM with memory type as specified in [A, T]MTypeDram and data accesses are directed at MMIO space and with attributes based on [A, T]MTypeIoWc.
  - If not in SMM, then the accesses are directed at MMIO space with attributes based on [A,T]MTypeIoWc.




# X86 goes to Harvard









# X86 goes to Harvard



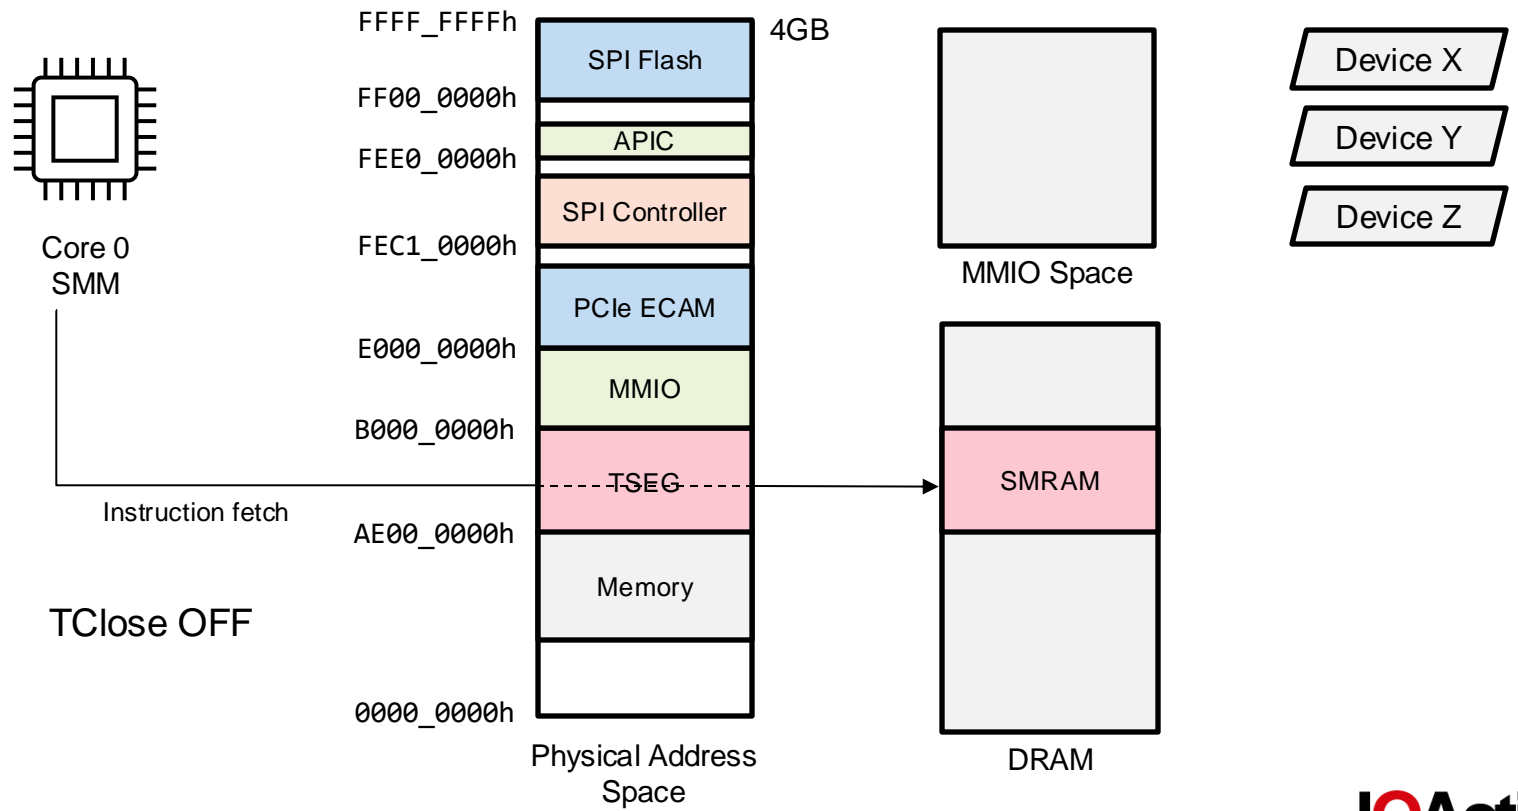
Core 0  
SMM

TClose OFF





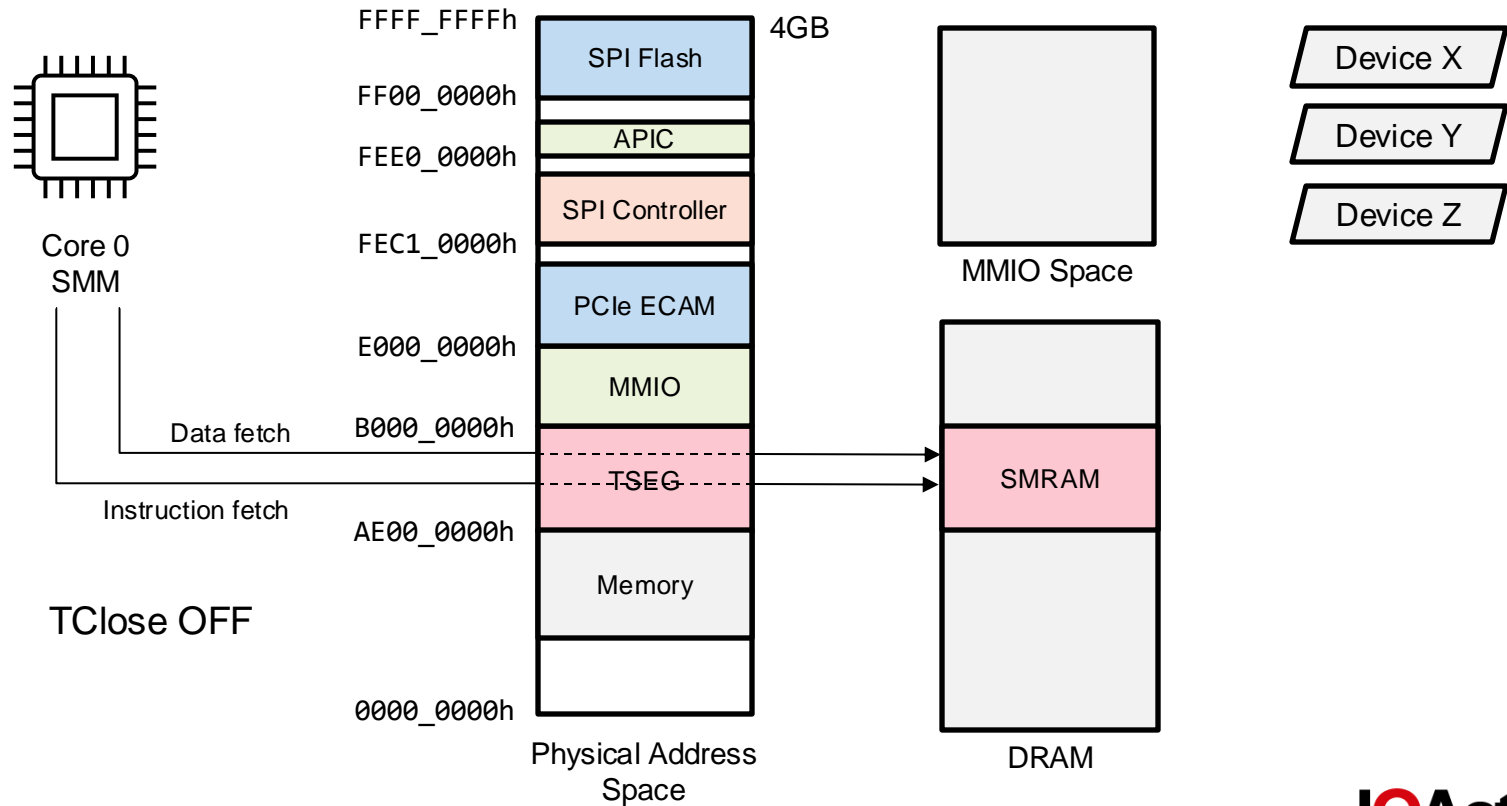
# X86 goes to Harvard








# X86 goes to Harvard






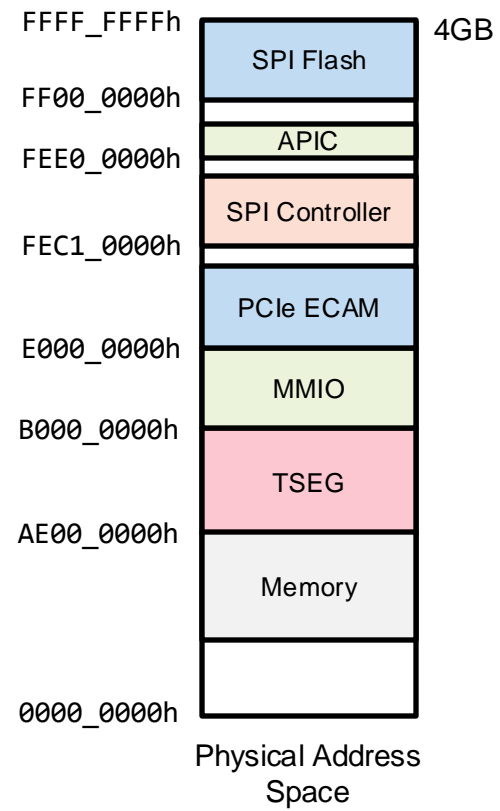


# X86 goes to Harvard



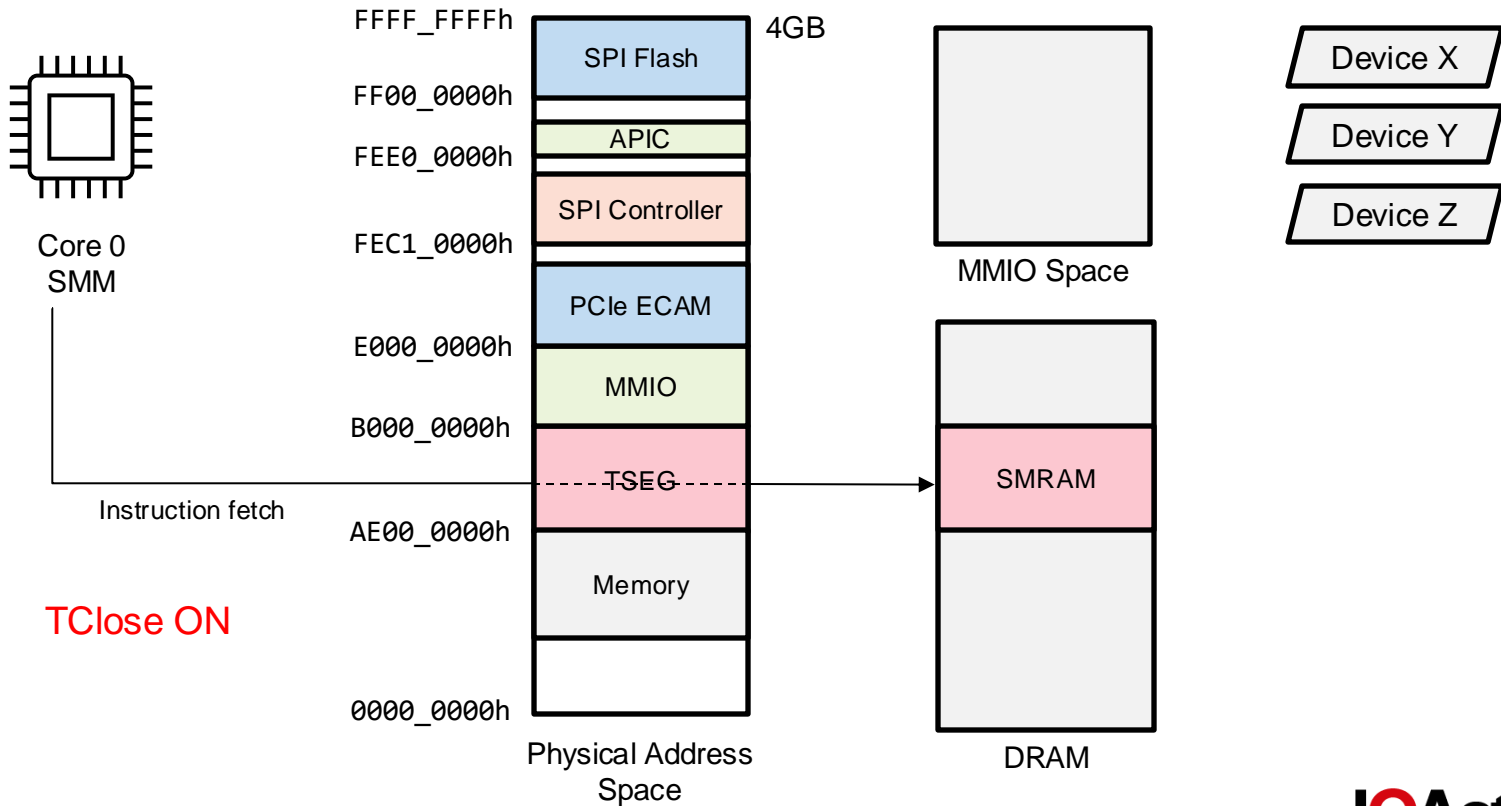
Core 0  
SMM

TClose ON



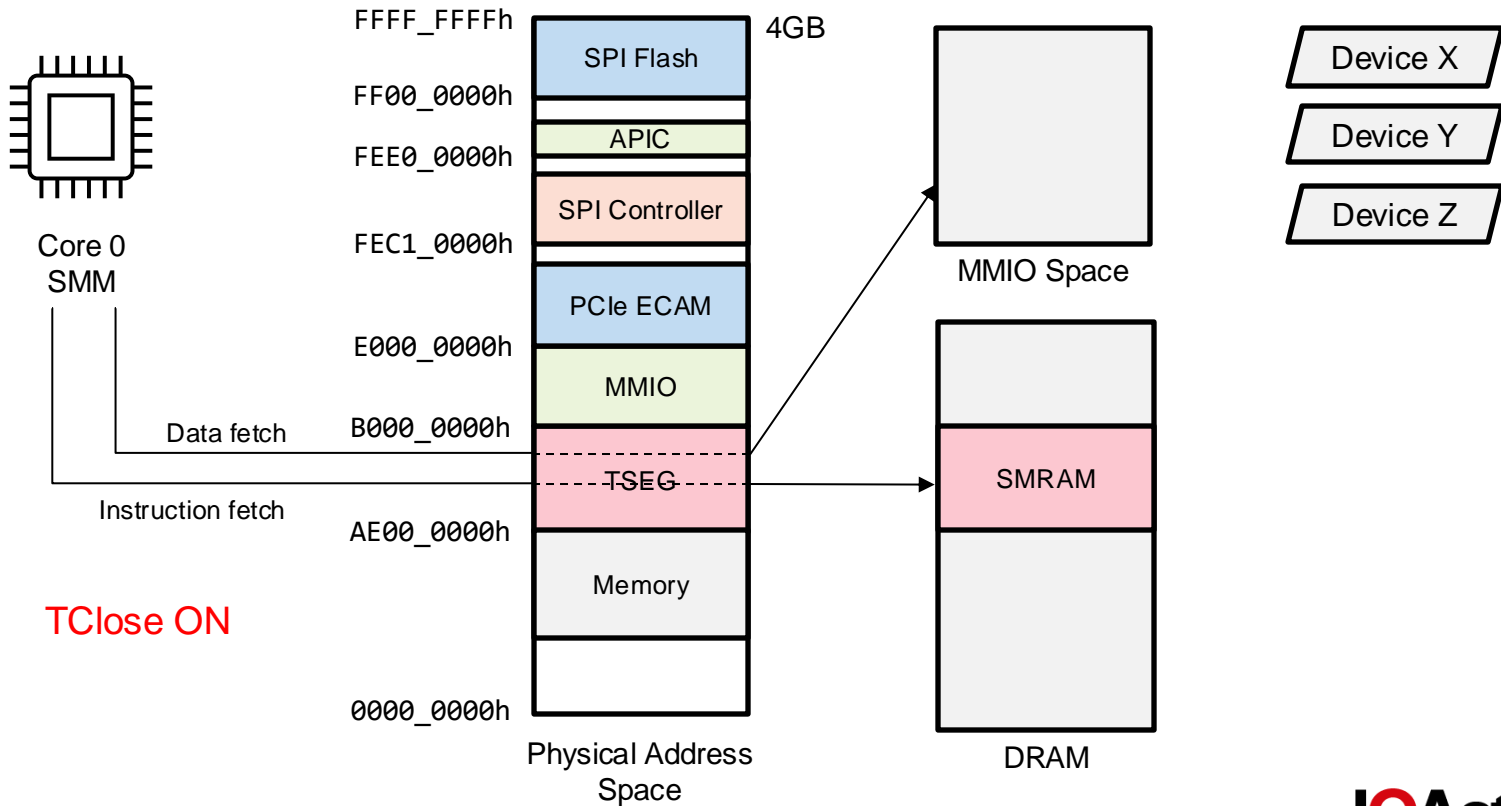


# X86 goes to Harvard



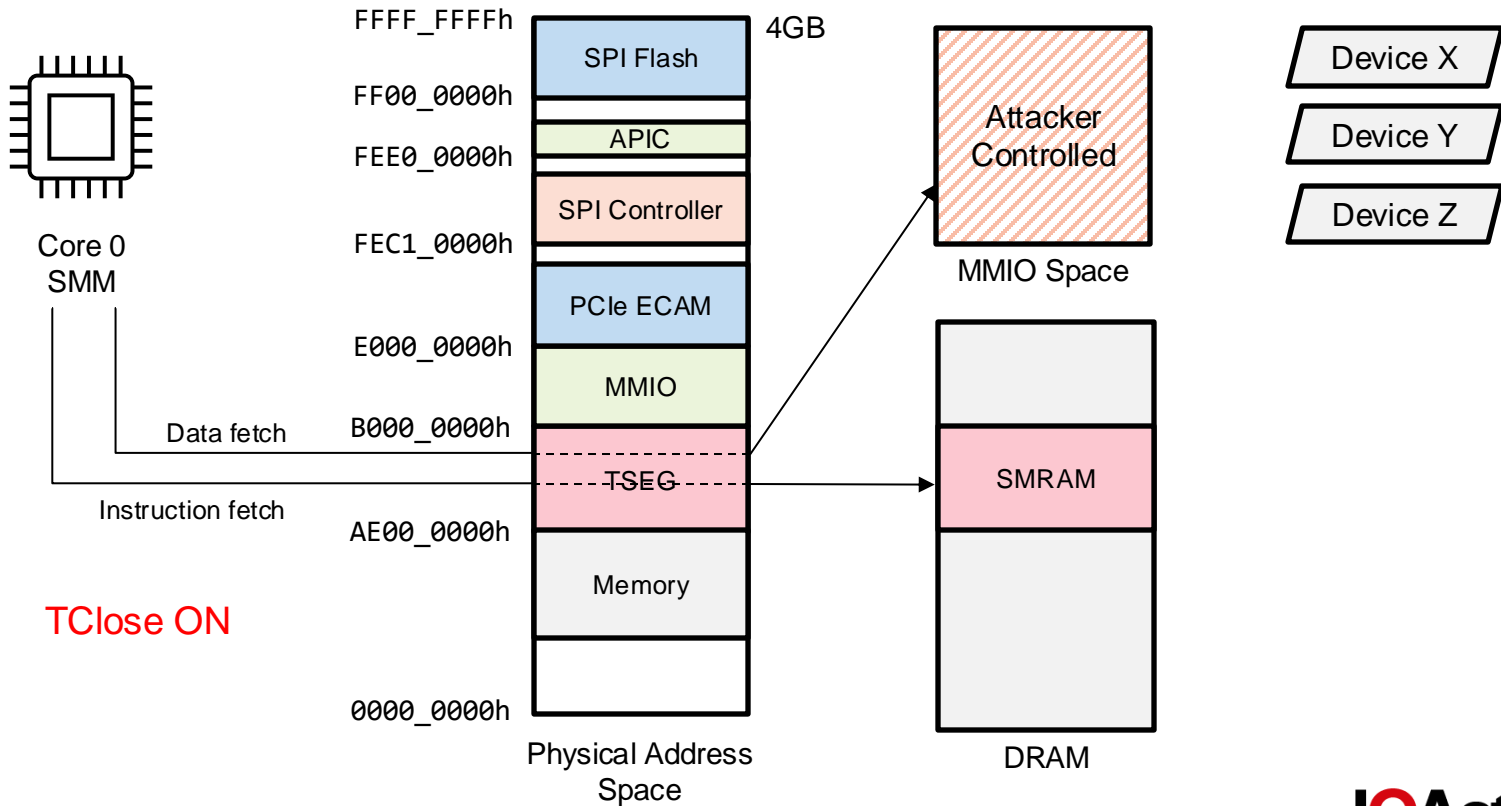


# X86 goes to Harvard





# X86 goes to Harvard





# Triggering the condition

```
void test() {  
    open_platbox_device();  
  
    UINT64 tseg_mask = 0;  
    do_read_msr(AMD_MSR_SMM_TSEG_MASK, &tseg_mask);  
    tseg_mask = tseg_mask | (0b11 << 2);  
    do_write_msr(AMD_MSR_SMM_TSEG_MASK, tseg_mask);  
  
    SW_SMI_CALL smi_call = { 0 };  
    trigger_smi(&smi_call);  
  
    close_platbox_device();  
}
```



# Why does this feature exist?

- This allows to re-use the physical address space
- We have yet to see a vendor using this feature

## 8.11.5 Closing SMM

Sometimes within SMM code with ASeg or TSeg enabled, there is a requirement to access the I/O space at the same address as the current SMM segment. That is typically only accessible outside of SMM. To accomplish this function, the Aclose and Tclose bits from SMM\_MASK register are used. When the Aclose bit is set, data cache accesses to the ASeg that would normally go to DRAM are redirected to I/O, with the memory type specified by AMTypeIoWc.

The same function applies to the TSeg. Instruction cache accesses and Page Directory/Table accesses still access the SMM code in DRAM. When the SMM handler is done accessing the I/O space, it must clear the appropriate close bit. Failure to do so and then issuing an RSM will probably cause the processor to enter shutdown, as the save state is read from I/O space.



# When did this feature appear?

- First mentioned for AMD 0Fh processor families (2006)
- BIOS and Kernel Developer's Guide for AMD NPT Family 0Fh Processors  
<https://www.amd.com/content/dam/amd/en/documents/archived-tech-docs/programmer-references/32559.pdf>
- It's been around for 18 years...





# Differences with the "Memory Sinkhole"

- Cristopher Domas presented the Memory Sinkhole attack in 2015
  - Affected Intel Sandy Bridge and previous generations
  - Remaps the APIC over the TSEG area
  - Causes data fetches to go to MMIO instead of SMRAM
- Key differences:
  - The memory sinkhole only affects the 4K portion where the APIC gets mapped
  - Sinkclose changes the behavior of the entire TSEG region
    - Any device could be overlapped... right?



# Brainstorming attack ideas



# Attack idea

- Use a PCIe device with a BAR having register values such that when overlapped with the SMM entry point, we could take control of the execution
- There are multiple integrated devices in modern systems
- We can try re-mapping the PCI Base Address Register (BAR) from one of them to make it overlap with SMRAM
- The registers for the device should become visible for the OS at the TSEG location



# PCI BARs failed

```
/dev/KernetixDriver0 opened successfully: 3
+ SMM region info:
TSEG Base   : bf000000
TSEG Size   : 00ffffff
SMM Base    : bfea8000
SMM-Entry   : bfeb0000
Ethernet controller BAR2 at: d0714000
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
-> remapping BAR2 to overlap TSEG
+ successfully overlapped the ethernet bar over SMM at: bfeb0000
-> view of memory at smm entry point:
0xbfeb0000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

-> Memory at BAR2 (d0714000):
0xd0714000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

Restoring BAR and dumping again:
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
```



# PCI BARs failed

```
/dev/KernetixDriver0 opened successfully: 3
+ SMM region info:
TSEG Base   : bf000000
TSEG Size   : 00ffffff
SMM Base    : bfea8000
SMM-Entry   : bfeb0000
Ethernet controller BAR2 at: d0714000
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
-> remapping BAR2 to overlap TSEG
+ successfully overlapped the ethernet bar over SMM at: bfeb0000
-> view of memory at smm entry point:
0xbfeb0000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

-> Memory at BAR2 (d0714000):
0xd0714000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

Restoring BAR and dumping again:
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
```

Visible device registers



# PCI BARs failed

```
/dev/KernetixDriver0 opened successfully: 3
+ SMM region info:
TSEG Base   : bf000000
TSEG Size   : 00ffffff
SMM Base    : bfea8000
SMM-Entry   : bfeb0000
Ethernet controller BAR2 at: d0714000
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
-> remapping BAR2 to overlap TSEG
+ successfully overlapped the ethernet bar over SMM at: bfeb0000
-> view of memory at smm entry point:
0xbfeb0000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

-> Memory at BAR2 (d0714000):
0xd0714000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

Restoring BAR and dumping again:
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
```

Visible device registers

Remap failed; registers are not available



# PCI BARs failed

```
/dev/KernetixDriver0 opened successfully: 3
+ SMM region info:
TSEG Base   : bf000000
TSEG Size   : 00ffffff
SMM Base    : bfea8000
SMM-Entry   : bfeb0000
Ethernet controller BAR2 at: d0714000
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
-> remapping BAR2 to overlap TSEG
+ successfully overlapped the ethernet bar over SMM at: bfeb0000
-> view of memory at smm entry point:
0xbfeb0000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
-> Memory at BAR2 (d0714000):
0xd0714000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
Restoring BAR and dumping again:
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
```

Visible device registers

Remap failed; registers are not available

The BAR was indeed moved from its original place



# PCI BARs failed

```
/dev/KernetixDriver0 opened successfully: 3
+ SMM region info:
TSEG Base   : bf000000
TSEG Size   : 00ffffff
SMM Base    : bfea8000
SMM-Entry   : bfeb0000
Ethernet controller BAR2 at: d0714000
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
-> remapping BAR2 to overlap TSEG
+ successfully overlapped the ethernet bar over SMM at: bfeb0000
-> view of memory at smm entry point:
0xbfeb0000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xbfeb0020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

-> Memory at BAR2 (d0714000):
0xd0714000 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714010 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....
0xd0714020 | ff ff ff ff ff ff ff ff ff ff ff ff ff ff ff | .....

Restoring BAR and dumping again:
0xd0714000 | 00 2b 67 52 7c c0 00 00 40 00 00 00 80 00 00 00 | .+gR|...@.....
0xd0714010 | 00 c0 ff ff 00 00 00 00 08 07 06 00 00 00 00 00 | .....
0xd0714020 | 00 b0 ff ff 00 00 00 00 00 00 00 00 00 00 00 00 | .....
```

Visible device registers

Remap failed; registers are not available

The BAR was indeed moved from its original place

After restoration





# TOM - Top of Memory

## MSRC001\_001A Top Of Memory (TOP\_MEM)


Reset: 0000\_0000\_0000\_0000h.

Bits	Description
63:40	RAZ.
39:23	<b>TOM[39:23]: top of memory.</b> Read-write. Specifies the address that divides between MMIO and DRAM. This value is normally placed below 4G. From TOM to 4G is MMIO; below TOM is DRAM. See <a href="#">2.4.6 [System Address Map]</a> and <a href="#">2.9.11 [DRAM CC6/PC6 Storage]</a> .
22:0	RAZ.

- This register dictates where the MMIO region below 4G starts
- On Intel this register has a lock bit and cannot be modified when set
- There is no such lock in AMD :)





# Moving TOM down





# Moving TOM down





# Moving TOM down



*This worked in theory but not in practice...*



# Memory routing priorities

## 2.4.6.1.2 Determining The Access Destination for Core Accesses

The access destination, DRAM or MMIO, is based on the highest priority of the following ranges that the access falls in: 1==Lowest priority.

1. RdDram/WrDram as determined by [MSRC001\\_001A \[Top Of Memory \(TOP\\_MEM\)\]](#) and [MSRC001\\_001D \[Top Of Memory 2 \(TOM2\)\]](#).
2. The IORRs. (see [MSRC001\\_00\[18,16\]](#) and [MSRC001\\_00\[19,17\]](#)).
3. The fixed MTRRs. (see [MSR0000\\_02\[6F:68,59:58,50\]](#) [Fixed-Size MTRRs])
4. TSeg & ASeg SMM mechanism. (see [MSRC001\\_0112](#) and [MSRC001\\_0113](#))
5. MMIO config space, APIC space.
  - MMIO APIC space and MMIO config space must not overlap.
  - RdDram=IO, WrDram=IO.
  - See [2.4.9.1.2 \[APIC Register Space\]](#) and [2.7 \[Configuration Space\]](#).
6. NB address space routing. See [2.8.2.1.1 \[DRAM and MMIO Memory Space\]](#).



# Memory routing priorities

## 2.4.6.1.2 Determining The Access Destination for Core Accesses

The access destination, DRAM or MMIO, is based on the highest priority of the following ranges that the access falls in: 1==Lowest priority.

1. RdDram/WrDram as determined by MSRC001\_001A [Top Of Memory (TOP\_MEM)] and MSRC001\_001D [Top Of Memory 2 (TOM2)].
2. The IORRs. (see MSRC001\_00[18,16] and MSRC001\_00[19,17]).
3. The fixed MTRRs. (see MSR0000\_02[6F:68,59:58,50] [Fixed-Size MTRRs])
4. TSeg & ASeg SMM mechanism. (see MSRC001\_0112 and MSRC001\_0113)
5. MMIO config space, APIC space.
  - MMIO APIC space and MMIO config space must not overlap.
  - RdDram=IO, WrDram=IO.
  - See 2.4.9.1.2 [APIC Register Space] and 2.7 [Configuration Space].
6. NB address space routing. See 2.8.2.1.1 [DRAM and MMIO Memory Space].





# Memory routing priorities

## 2.4.6.1.2 Determining The Access Destination for Core Accesses

The access destination, DRAM or MMIO, is based on the highest priority of the following ranges that the access falls in: 1==Lowest priority.

- ~~1. RdDrAm/WrDrAm as determined by MSRC001\_001A [Top Of Memory (TOP\_MEM)] and MSRC001\_001D [Top Of Memory 2 (TOM2)].~~
- ~~2. The IORRs. (see MSRC001\_00[18,16] and MSRC001\_00[19,17])~~
- ~~3. The fixed MTRRs. (see MSR0000\_02[6F:68,59:58,50] [Fixed-Size MTRRs])~~
4. TSeg & ASeg SMM mechanism. (see MSRC001\_0112 and MSRC001\_0113)
5. MMIO config space, APIC space.
  - MMIO APIC space and MMIO config space must not overlap.
  - RdDrAm=IO, WrDrAm=IO.
  - See 2.4.9.1.2 [APIC Register Space] and 2.7 [Configuration Space].
6. NB address space routing. See 2.8.2.1.1 [DRAM and MMIO Memory Space].



# Memory routing priorities

## 2.4.6.1.2 Determining The Access Destination for Core Accesses

The access destination, DRAM or MMIO, is based on the highest priority of the following ranges that the access falls in: 1==Lowest priority.

- ~~1. RdDram/WrDram as determined by MSRC001\_001A [Top Of Memory (TOP\_MEM)] and MSRC001\_001D [Top Of Memory 2 (TOM2)].~~
- ~~2. The IORRs. (see MSRC001\_00[18,16] and MSRC001\_00[19,17])~~
- ~~3. The fixed MTRRs. (see MSR0000\_02[6F:68,59:58,50] [Fixed-Size MTRRs])~~
4. TSeg & ASeg SMM mechanism. (see MSRC001\_0112 and MSRC001\_0113)
5. MMIO config space, APIC space.
  - MMIO APIC space and MMIO config space must not overlap.
  - RdDram=IO, WrDram=IO.
  - See 2.4.9.1.2 [APIC Register Space] and 2.7 [Configuration Space].
6. NB address space routing. See 2.8.2.1.1 [DRAM and MMIO Memory Space].





# Memory routing priorities


## 2.4.6.1.2 Determining The Access Destination for Core Accesses

The access destination, DRAM or MMIO, is based on the highest priority of the following ranges that the access falls in: 1==Lowest priority.


- ~~1. RdDrAm/WrDrAm as determined by MSRC001\_001A [Top Of Memory (TOP\_MEM)] and MSRC001\_001D [Top Of Memory 2 (TOM2)].~~
- ~~2. The IORRs. (see MSRC001\_00[18,16] and MSRC001\_00[19,17])~~
- ~~3. The fixed MTRRs. (see MSR0000\_02[6F:68,59:58,50] [Fixed-Size MTRRs])~~
4. TSeg & ASeg SMM mechanism. (see MSRC001\_0112 and MSRC001\_0113)
5. MMIO config space, APIC space.
  - MMIO APIC space and MMIO config space must not overlap.
  - RdDrAm=IO, WrDrAm=IO.
  - See 2.4.9.1.2 [APIC Register Space] and 2.7 [Configuration Space].
6. NB address space routing. See 2.8.2.1.1 [DRAM and MMIO Memory Space]. ?



# Analysis of the SMM entry point



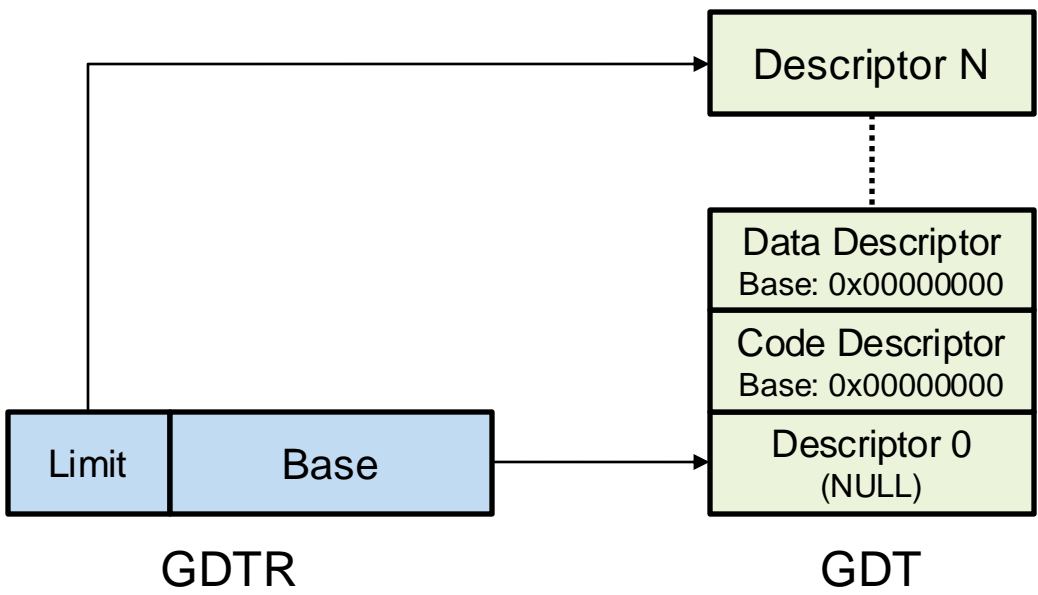
SMM  
Mode





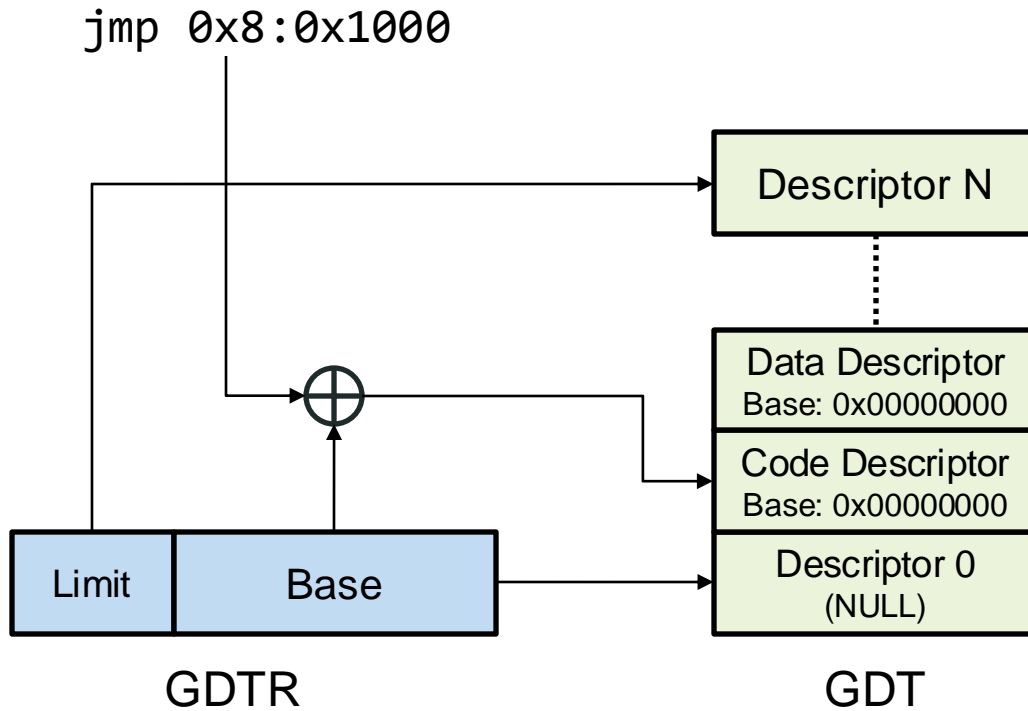
# Global Descriptor Table (GDT)

jmp 0x8:0x1000



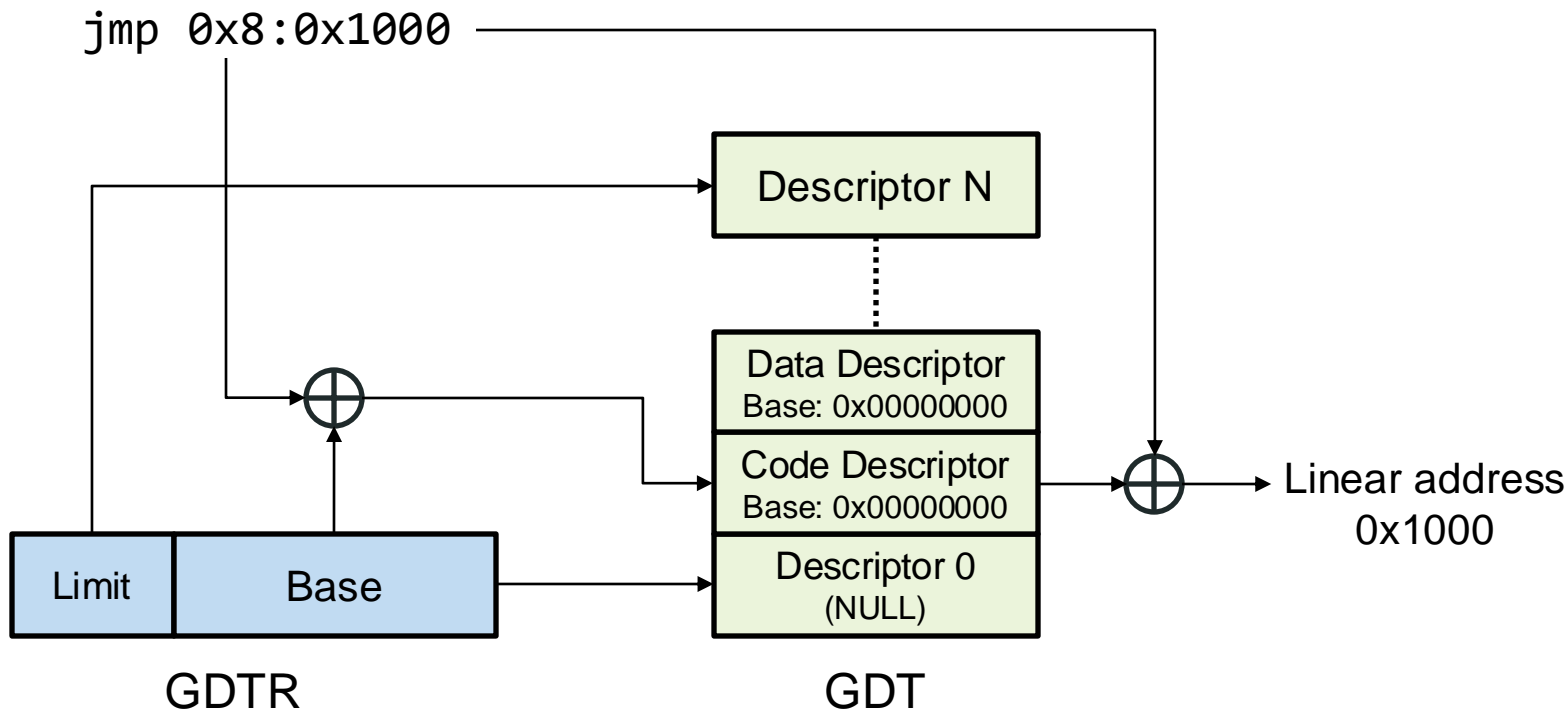


# Global Descriptor Table (GDT)





# Global Descriptor Table (GDT)





# Analysis of the EDKII SMM entry point

```
0:  bb 4d 80          mov     bx,0x804d      ; 0x8000 + 0x4D
3:  2e a1 d8 fd       mov     ax,cs:0xfdd8   ; DSC_OFFSET + 0xD8
7:  48                dec     ax
8:  2e 89 07          mov     WORD PTR cs:[bx],ax
b:  2e 66 a1 d0 fd     mov     eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
10: 2e 66 89 47 02     mov     DWORD PTR cs:[bx+0x2],eax
15: 2e 66 0f 01 17     lgdt    cs:[bx];
1a: b8 08 00          mov     ax,0x8
1d: 2e 89 47 fe       mov     WORD PTR cs:[bx-0x2],ax
21: 66 bf 00 30 f4 ae  mov     edi,0xae43000
27: 66 67 8d 87 53 80 00 lea     eax,[edi+0x8053]
2e: 00
2f: 2e 66 89 47 fa     mov     DWORD PTR cs:[bx-0x6],eax
34: 0f 20 c3          mov     ebx,cr0
37: 66 81 e3 f3 ff fa 9f and     ebx,0x9ffa0003
3e: 66 83 cb 23       or      ebx,0x23
42: 0f 22 c3          mov     cr0,ebx
45: 66 ea 53 b0 f4 ae 08 jmp     0x8:0xae4b053
4c: 00
4d: [ GDTR HERE ]
```



# Analysis of the EDKII SMM entry point

```
0:  bb 4d 80          mov     bx,0x804d      ; 0x8000 + 0x4D → SMM entry point + 0x4D
3:  2e a1 d8 fd       mov     ax,cs:0xfdd8   ; DSC_OFFSET + 0xD8
7:  48                dec     ax
8:  2e 89 07          mov     WORD PTR cs:[bx],ax
b:  2e 66 a1 d0 fd     mov     eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
10: 2e 66 89 47 02     mov     DWORD PTR cs:[bx+0x2],eax
15: 2e 66 0f 01 17     lgdt    cs:[bx];
1a: b8 08 00          mov     ax,0x8
1d: 2e 89 47 fe       mov     WORD PTR cs:[bx-0x2],ax
21: 66 bf 00 30 f4 ae  mov     edi,0xae43000
27: 66 67 8d 87 53 80 00 lea     eax,[edi+0x8053]
2e: 00
2f: 2e 66 89 47 fa     mov     DWORD PTR cs:[bx-0x6],eax
34: 0f 20 c3          mov     ebx,cr0
37: 66 81 e3 f3 ff fa 9f and     ebx,0x9ffa3fff
3e: 66 83 cb 23       or      ebx,0x23
42: 0f 22 c3          mov     cr0,ebx
45: 66 ea 53 b0 f4 ae 08 jmp     0x8:0xae4b053
4c: 00
4d: [ GDTR HERE ]
```



# Analysis of the EDKII SMM entry point

```
0:  bb 4d 80          mov     bx,0x804d      ; 0x8000 + 0x4D → SMM entry point + 0x4D
3:  2e a1 d8 fd       mov     ax,cs:0xfdd8   ; DSC_OFFSET + 0xD8
7:  48                dec     ax
8:  2e 89 07          mov     WORD PTR cs:[bx],ax
b:  2e 66 a1 d0 fd     mov     eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
10: 2e 66 89 47 02     mov     DWORD PTR cs:[bx+0x2],eax
15: 2e 66 0f 01 17     lgdt    cs:[bx];
1a: b8 08 00          mov     ax,0x8
1d: 2e 89 47 fe       mov     WORD PTR cs:[bx-0x2],ax
21: 66 bf 00 30 f4 ae   mov     edi,0xae43000
27: 66 67 8d 87 53 80 00 lea     eax,[edi+0x8053]
2e: 00
2f: 2e 66 89 47 fa     mov     DWORD PTR cs:[bx-0x6],eax
34: 0f 20 c3          mov     ebx,cr0
37: 66 81 e3 f3 ff fa 9f and     ebx,0x9ffaaff3
3e: 66 83 cb 23       or      ebx,0x23
42: 0f 22 c3          mov     cr0,ebx
45: 66 ea 53 b0 f4 ae 08 jmp     0x8:0xae4b053
4c: 00
4d:  [ GDTR HERE ]
```





# Analysis of the EDKII SMM entry point

```
0:  bb 4d 80          mov     bx,0x804d      ; 0x8000 + 0x4D → SMM entry point + 0x4D
3:  2e a1 d8 fd       mov     ax,cs:0xfdd8   ; DSC_OFFSET + 0xD8
7:  48                dec     ax
8:  2e 89 07          mov     WORD PTR cs:[bx],ax
b:  2e 66 a1 d0 fd     mov     eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
10: 2e 66 89 47 02     mov     DWORD PTR cs:[bx+0x2],eax
15: 2e 66 0f 01 17     lgdt    cs:[bx]; → Loads GDTR
1a: b8 08 00          mov     ax,0x8
1d: 2e 89 47 fe       mov     WORD PTR cs:[bx-0x2],ax
21: 66 bf 00 30 f4 ae  mov     edi,0xae43000
27: 66 67 8d 87 53 80 00 lea     eax,[edi+0x8053]
2e: 00
2f: 2e 66 89 47 fa     mov     DWORD PTR cs:[bx-0x6],eax
34: 0f 20 c3          mov     ebx,cr0
37: 66 81 e3 f3 ff fa 9f and     ebx,0x9ffa3fff
3e: 66 83 cb 23       or      ebx,0x23
42: 0f 22 c3          mov     cr0,ebx
45: 66 ea 53 b0 f4 ae 08 jmp     0x8:0xae4b053
4c: 00
4d: [ GDTR HERE ]
```



# Analysis of the EDKII SMM entry point

```
0:  bb 4d 80          mov     bx,0x804d      ; 0x8000 + 0x4D → SMM entry point + 0x4D
3:  2e a1 d8 fd       mov     ax,cs:0xfdd8   ; DSC_OFFSET + 0xD8
7:  48                dec     ax
8:  2e 89 07          mov     WORD PTR cs:[bx],ax
b:  2e 66 a1 d0 fd     mov     eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
10: 2e 66 89 47 02     mov     DWORD PTR cs:[bx+0x2],eax
15: 2e 66 0f 01 17     lgdt    cs:[bx]; → Loads GDTR
1a: b8 08 00          mov     ax,0x8
1d: 2e 89 47 fe       mov     WORD PTR cs:[bx-0x2],ax
21: 66 bf 00 30 f4 ae   mov     edi,0xae43000
27: 66 67 8d 87 53 80 00 lea     eax,[edi+0x8053]
2e: 00
2f: 2e 66 89 47 fa     mov     DWORD PTR cs:[bx-0x6],eax
34: 0f 20 c3          mov     ebx,cr0
37: 66 81 e3 f3 ff fa 9f and     ebx,0x9ffa0003
3e: 66 83 cb 23       or      ebx,0x23
42: 0f 22 c3          mov     cr0,ebx
45: 66 ea 53 b0 f4 ae 08 jmp     0x8:0xae4b053 → Jumps to 32-bit (protected) code
4c: 00
4d:  [ GDTR HERE ]
```



# Analysis of the EDKII SMM entry point

```
0:  bb 4d 80          mov     bx,0x804d      ; 0x8000 + 0x4D → SMM entry point + 0x4D
3:  2e a1 d8 fd       mov     ax,cs:0xfdd8   ; DSC_OFFSET + 0xD8
7:  48                dec     ax
8:  2e 89 07          mov     WORD PTR cs:[bx],ax
b:  2e 66 a1 d0 fd     mov     eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
10: 2e 66 89 47 02     mov     DWORD PTR cs:[bx+0x2],eax
15: 2e 66 0f 01 17     lgdt    cs:[bx]; → Loads GDTR
1a: b8 08 00          mov     ax,0x8
1d: 2e 89 47 fe       mov     WORD PTR cs:[bx-0x2],ax
21: 66 bf 00 30 f4 ae  mov     edi,0xae43000
27: 66 67 8d 87 53 80 00 lea     eax,[edi+0x8053]
2e: 00
2f: 2e 66 89 47 fa     mov     DWORD PTR cs:[bx-0x6],eax
34: 0f 20 c3          mov     ebx,cr0
37: 66 81 e3 f3 ff fa 9f and     ebx,0x9ffa3fff
3e: 66 83 cb 23       or      ebx,0x23
42: 0f 22 c3          mov     cr0,ebx
45: 66 ea 53 b0 f4 ae 08 jmp     0x8:0xae4b053 → Jumps to 32-bit (protected) code
4c: 00
4d: [ GDTR HERE ]
```

*We need to control the BAR of the overlapped device at offset 0x4D*



# Problems with the APIC

- The system becomes unstable when the APIC is moved
- The APIC registers are not useful for taking control at the SMM entry point



## Writes are discarded

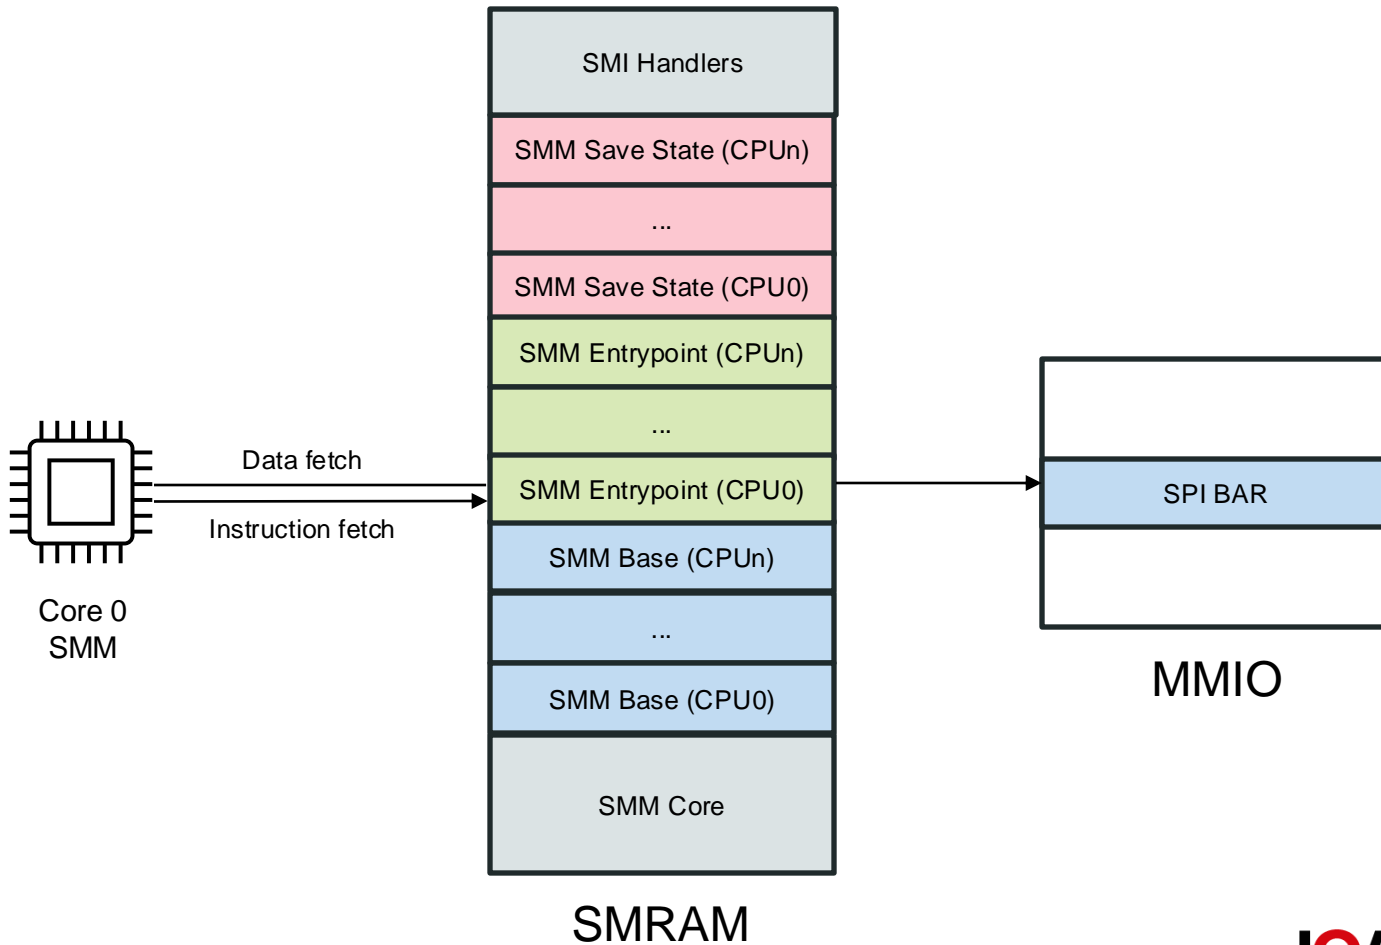


# Introducing the SPI controller



# SPI controller

- Used to read / write / erase the SPI flash
- Key features:
  - The BAR can be relocated over the SMM entry point
  - Portions of the BAR are attacker-controlled
  - Takes precedence over SMRAM when TClose is enabled





Memory

Address = 00000000FEC11000

Info Text

	00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F
00	00	00	20	0F	00	00	00	00	00	00	00	00	00	00	22	02
10	06	20	04	04	06	04	9F	05	03	0B	0A	02	FF	9A	00	3B
20	12	07	33	31	08	20	20	20	0C	14	06	0E	C0	D4	00	80
30	C0	14	08	46	03	00	00	00	FC	FC	FC	FC	FC	88	00	00
40	3B	6B	BB	EB	00	00	00	00	00	00	00	00	42	00	12	00
50	00	12	13	0C	3C	6C	BC	EC	08	46	00	00	00	00	00	00
60	00	00	00	00	FD	00	00	00	00	00	00	00	00	00	00	00
70	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
80	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
90	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
C0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
D0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
E0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
F0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Hardware



# SPI BAR

- GDTR is loaded from offset 0x4D
- Controllable fields:
  - 0x4C-50: FCH::LPCPCICFG::memoryrange
  - 0x50-54: FCH::LPCPCICFG::rom\_protect\_0




# Debugging setup



# Debugging Setup

- BAR buffer
  - PCI Squirrel with PCILeech firmware
  - Used for persistent memory across boot cycles
- SMM backdoor
  - Used for modifying code in SMM on-demand



PCIe Squirrel

Power supply


M2 to PCI 4x  
adapter



# Exploitation

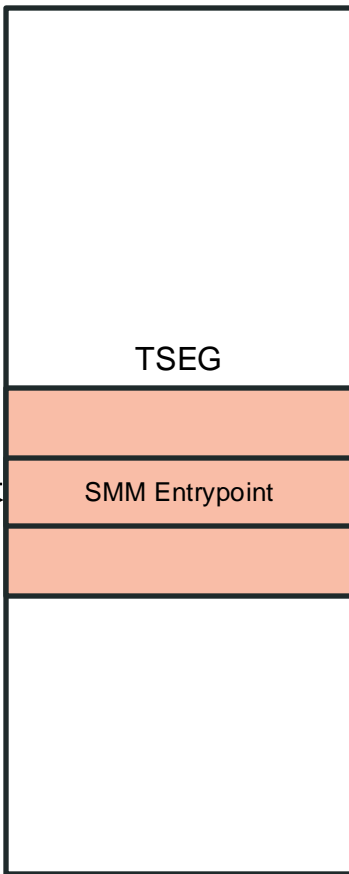


# Attempt #1



Core 0  
SMM

Data fetch  
Instruction fetch




DRAM

1. Remap


```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae4f3000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffaffff
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4f4b053
```




MMIO





Core 0  
SMM

Data fetch  
Instruction fetch




DRAM

1. Remap


```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffaffff
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```



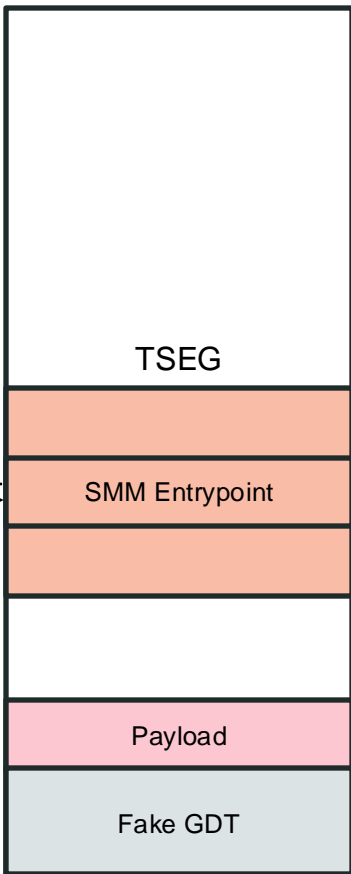
MMIO

2. Tweak



Core 0  
SMM

Data fetch  
Instruction fetch




DRAM

1. Remap

3. Map +  
tweak


```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```



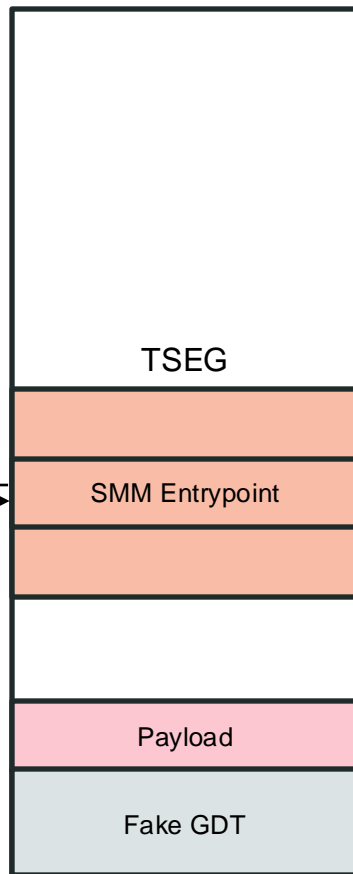
2. Tweak

MMIO



Core 0  
SMM

Data fetch  
Instruction fetch




DRAM

1. Remap

3. Map +  
tweak

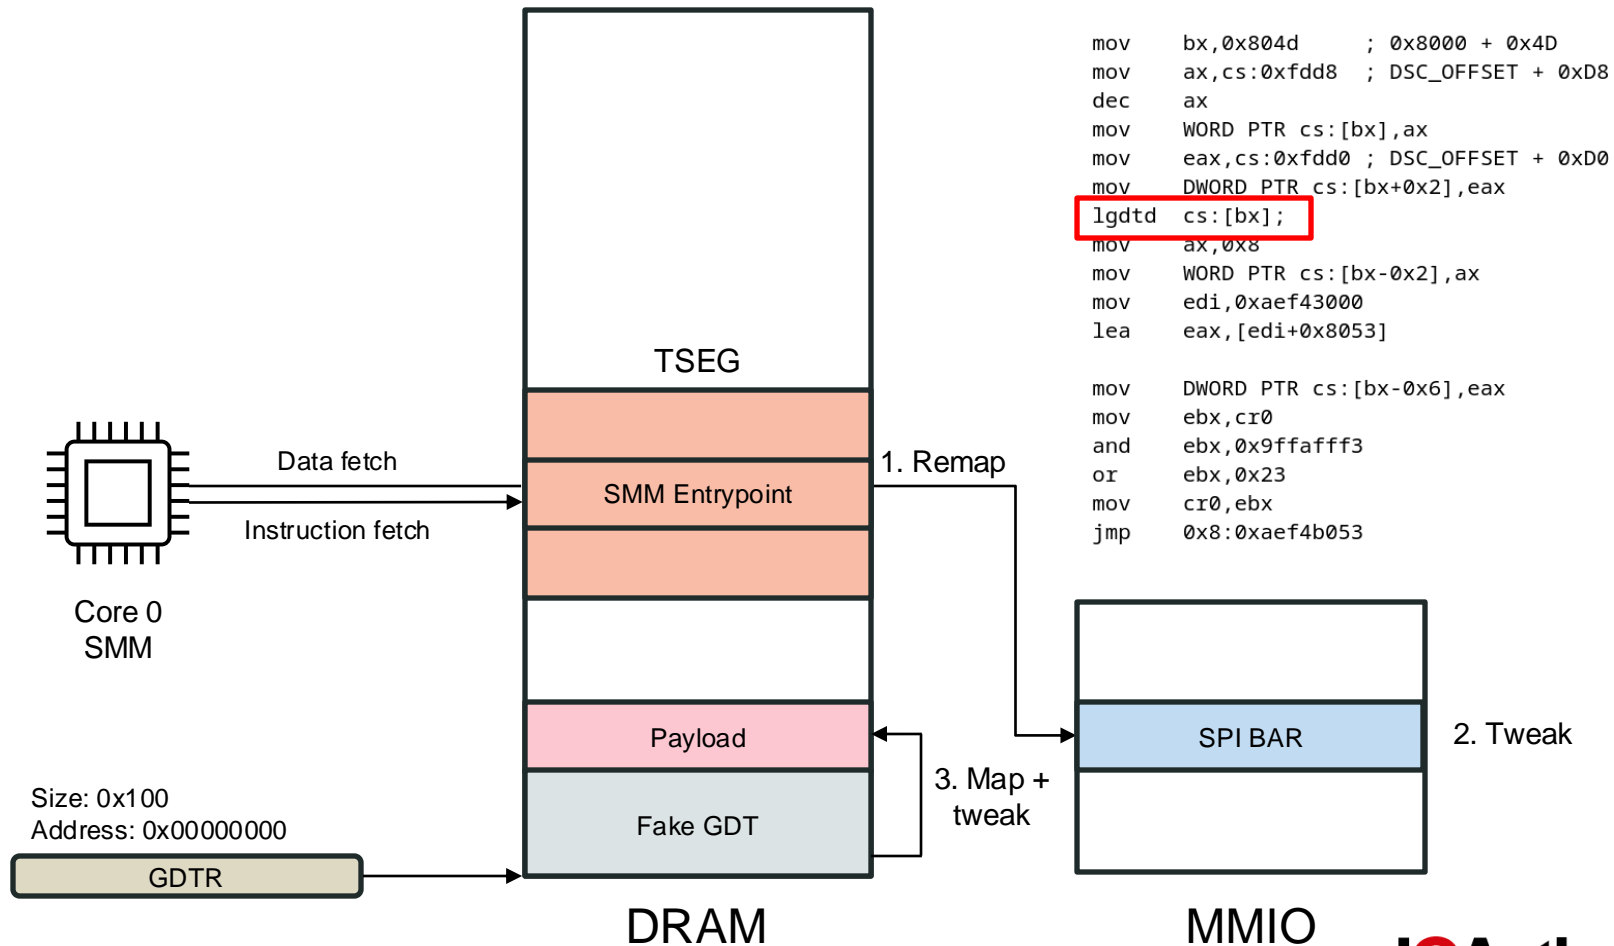
```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae4f3000
lea    eax,[edi+0x8053]

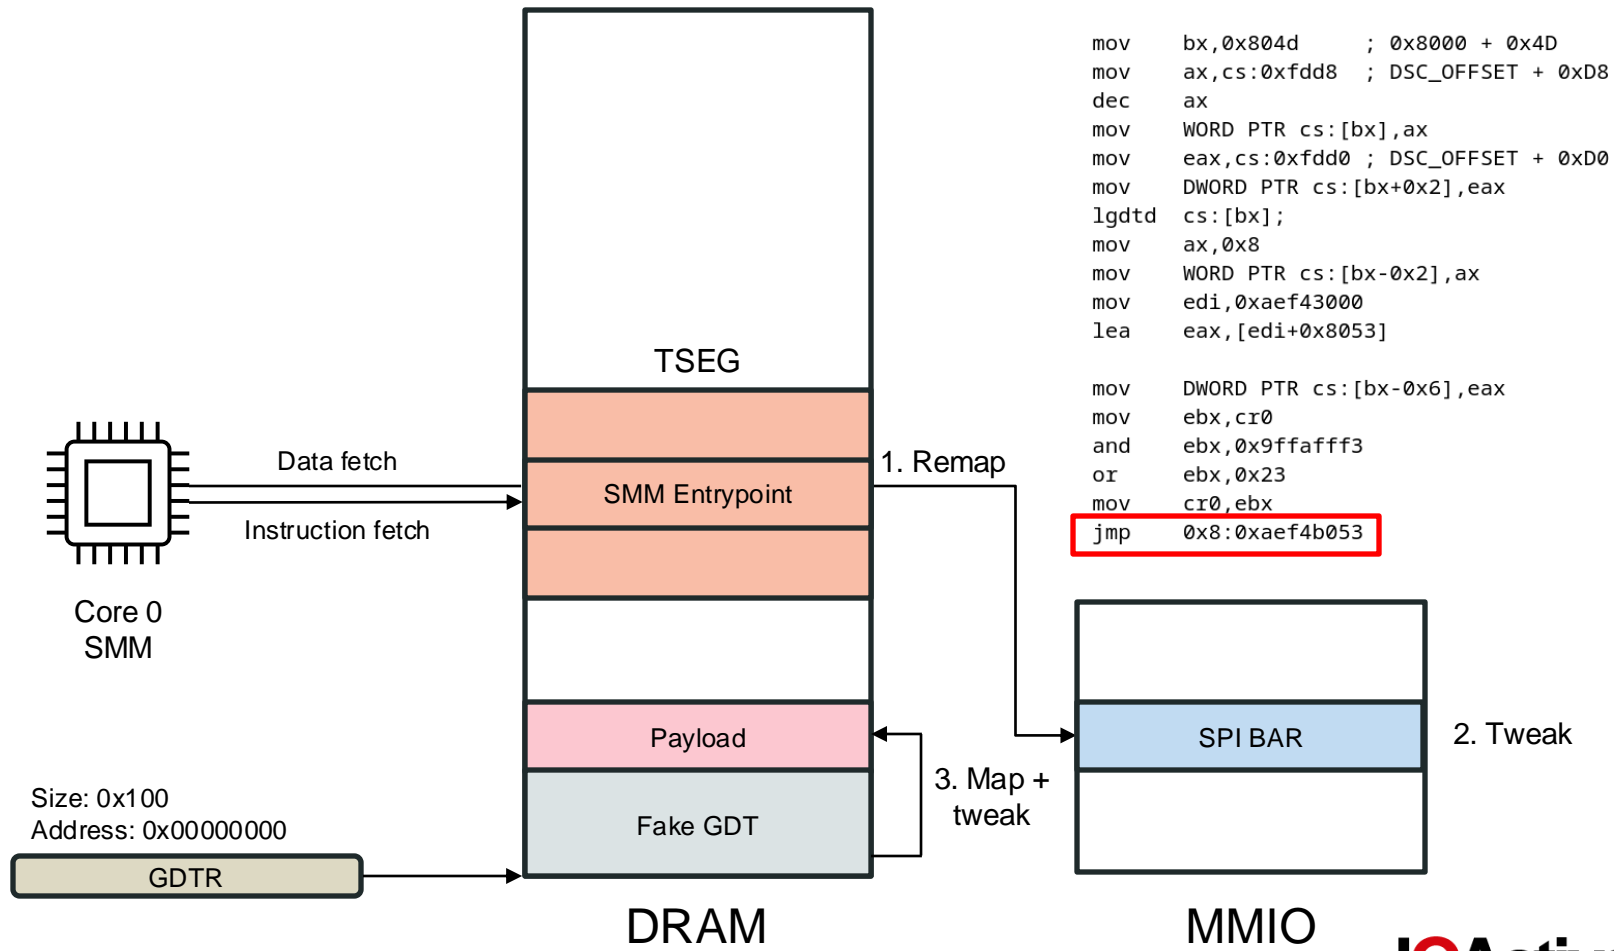
mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffaffff
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4f4b053
```

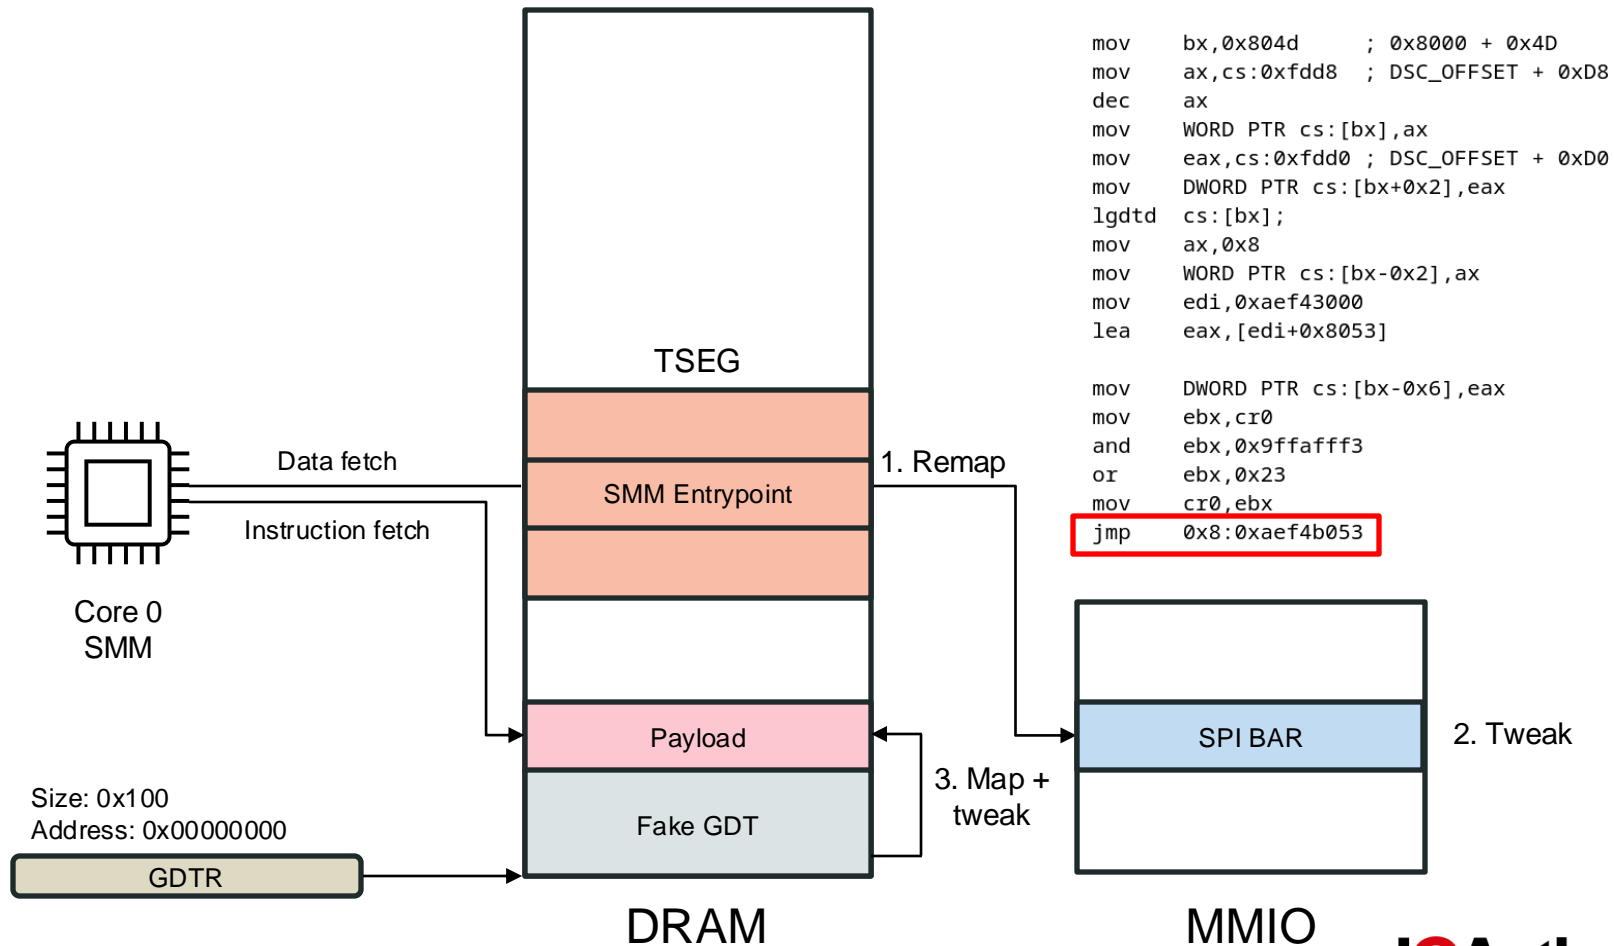


2. Tweak

MMIO

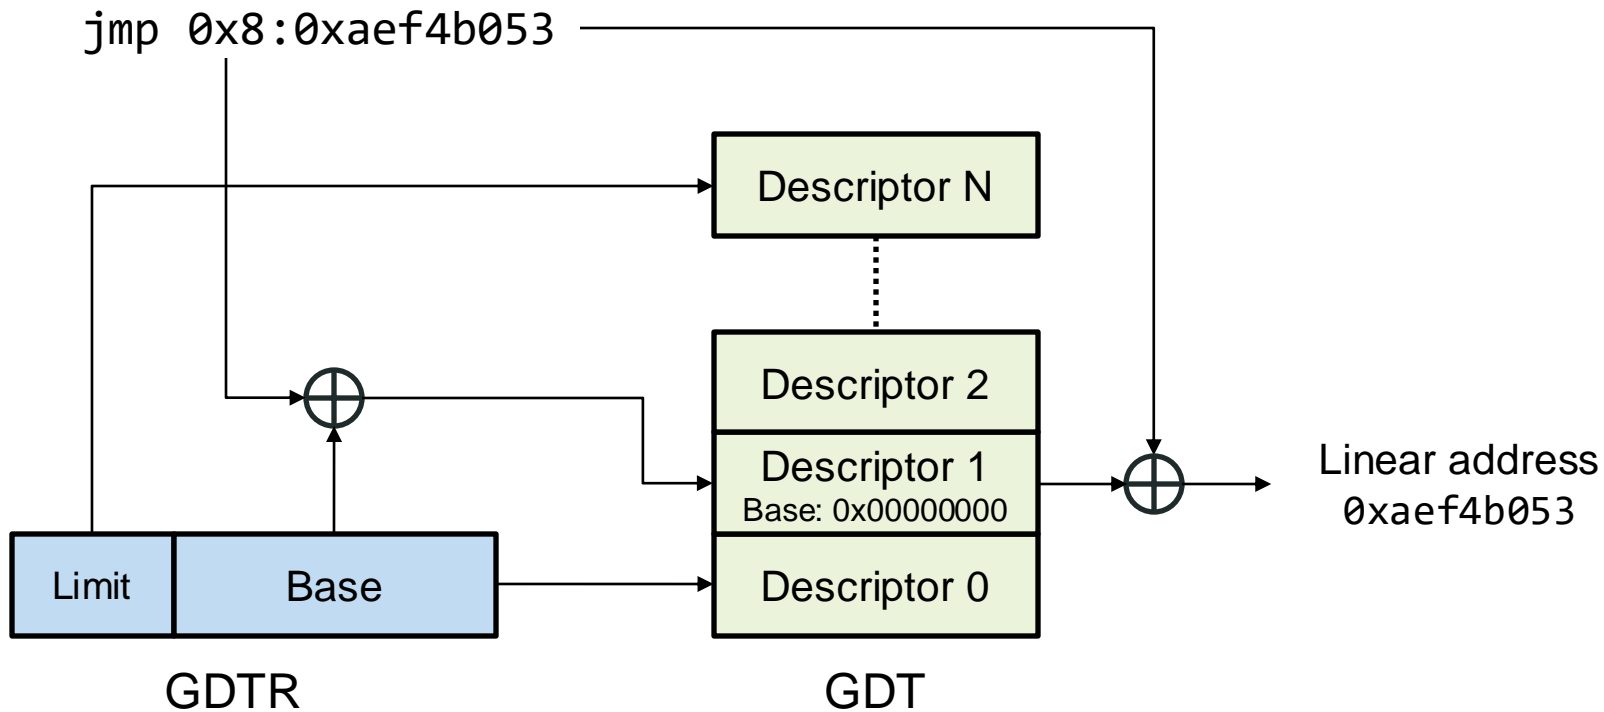






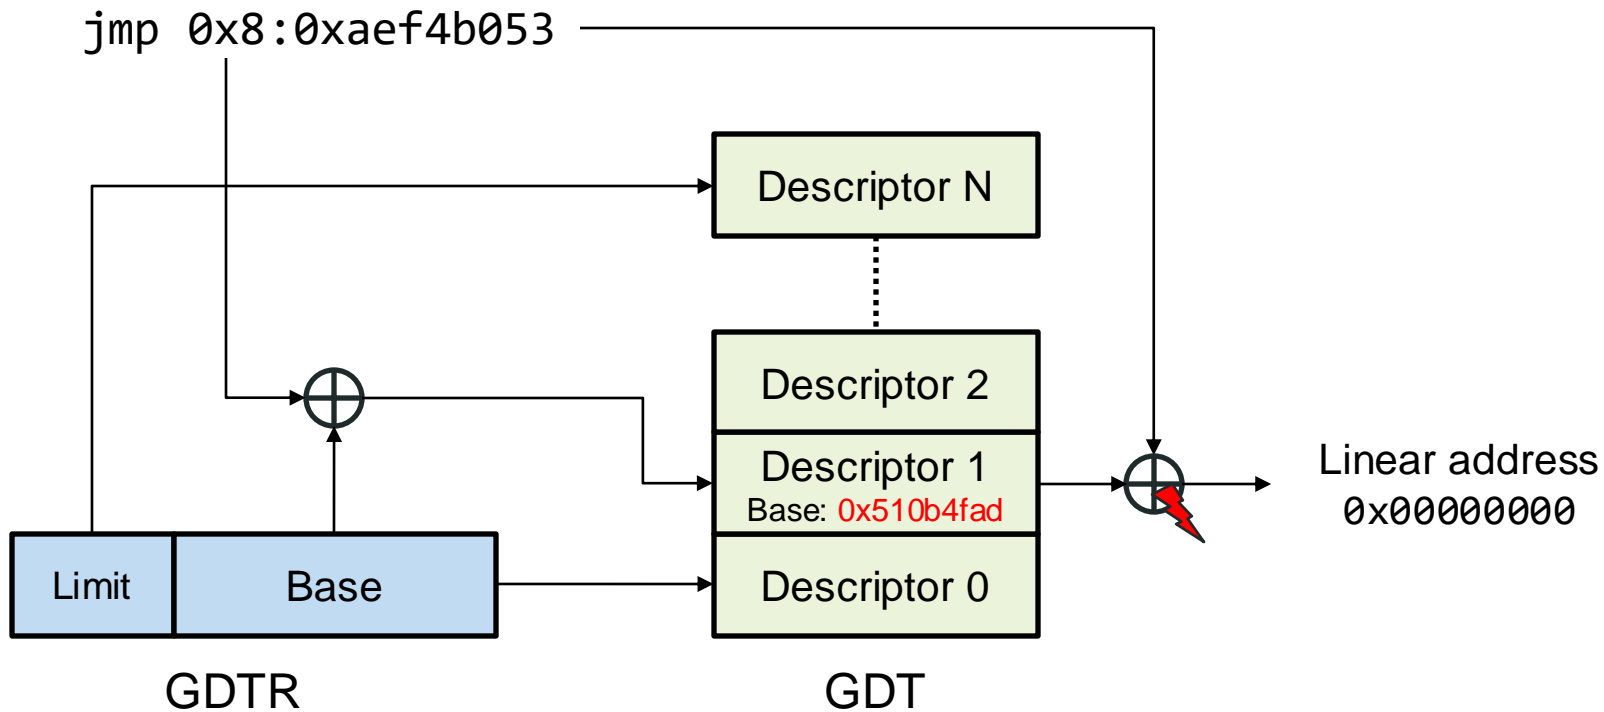


# GDT far jmp wrap-around





# GDT far jmp wrap-around







It worked, but the system crashed... why?



# The SMM save state

- The SMM save state is automatically saved upon entering SMM and restored when leaving it
  - With TClose enabled these writes are dropped
  - The SMM save state from the last SMI is still there
- Solution: Trigger SMI twice
  - Once without TClose to prime SMM save state
  - Once with TClose to trigger bug
- Does not require overwriting SMM save state values



# Attempt #2



The system crashed again... why?

# Enabling TClose

RW - Read & Write Utility v1.7 - [CPU MSR Registers]

Access Specific Window Help

index index index index SIO SPD I2C USB sm bios 55AA MPS E820 EDID ?

EC H90.CO 010 USB ?

Refresh

MTRR User

Register Name	Address	CPU1	CPU2	CPU3	CPU4	CPU5	CPU6
MTRR_DEF_TYPE	0x2FF	00000000000000C00	00000000000000C00	00000000000000C00	00000000000000C00	00000000000000C00	00000000000000C00
SMM_BASE	0xC0010111	000000000CEF38000	000000000CEF3A000	000000000CEF3C000	000000000CEF3E000	000000000CEF40000	000000000CEF42000
SMM_MASK	0xC0010113	0000FFFFFF006603	0000FFFFFF006603	0000FFFFFF006603	0000FFFFFF006603	0000FFFFFF006603	0000FFFFFF006603

Edit CPU1 MSR 0xC0010113

63	62	61	60	59	58	57	56	55	54	53	52	51	50	49	48	47	46	45	44	43	42	41	40	39	38	37	36	35	34	33	32
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	
00								00								FF								FF							

☐ To all CPUs

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	1	1	0	0	1	1	0	0	0	0	0	1	1	1	
FF								00								66								0F							

Done Cancel

# Bingo...

RW - Read & Write Utility v1.7 - [CPU MSR Registers]

Access Specific Window Help

index index index index SIO SPD I2C MSF Immo ACPI  
EC H00.CD 010 USB sm bios 55AA MPS E820 EDID ?

Refresh

MTRR User

Register Name	Address	CPU1	CPU2	CPU3	CPU4	CPU5	CPU6
MTRR_DEF_TYPE	0x2FF	0000000000000C00	0000000000000C00	0000000000000C00	0000000000000C00	0000000000000C00	0000000000000C00
SMM_BASE	0xC0010111	00000000CEF38000	00000000CEF3A000	00000000CEF3C000	00000000CEF3E000	00000000CEF40000	00000000CEF42000
SMM_MASK	0xC0010113	0000FFFFFF00660F	0000FFFFFF00660F	0000FFFFFF006603	0000FFFFFF006603	0000FFFFFF006603	0000FFFFFF006603




# Symmetric Multi-Threading

- Physical cores are split into two logical cores (threads)
- Some resources are shared between logical cores
  - SMM base MSR is separate but
  - TSEG mask MSR is not
- Is it an issue if only one core goes into SMM at a time?




# SMLs explained




Normal mode

```
xor eax, eax  
xor eax, eax
```




Normal mode

```
xor eax, eax  
xor eax, eax
```



Normal mode

```
xor eax, eax  
xor eax, eax
```




Normal mode

```
xor eax, eax  
xor eax, eax
```






# SMLs explained




Normal mode

```
xor eax, eax  
xor eax, eax  
smi
```




Normal mode

```
xor eax, eax  
xor eax, eax
```



Normal mode

```
xor eax, eax  
xor eax, eax
```




Normal mode

```
xor eax, eax  
xor eax, eax
```




# SMLs explained




Normal mode

```
xor eax, eax  
xor eax, eax  
smi
```




Normal mode

```
xor eax, eax  
xor eax, eax
```



Normal mode

```
xor eax, eax  
xor eax, eax
```




Normal mode

```
xor eax, eax  
xor eax, eax
```




# SMMs explained



SMM mode

```
xor eax, eax  
xor eax, eax  
smi
```


```
mov bs, 0x804d  
mov ax, cs:0xfdd8  
...
```



SMM mode

```
xor eax, eax  
xor eax, eax
```


```
mov bs, 0x804d  
mov ax, cs:0xfdd8  
...
```



SMM mode

```
xor eax, eax  
xor eax, eax
```

```
mov bs, 0x804d  
mov ax, cs:0xfdd8  
...
```



SMM mode

```
xor eax, eax  
xor eax, eax
```

```
mov bs, 0x804d  
mov ax, cs:0xfdd8  
...
```



# SMIs explained

- We assumed that SMIs are local, but they are global
- Initially we thought that:
  - we could control exactly which core enters into SMM first
  - each core would later reach the rendezvous routine and
  - send Inter-Processor-Interrupts (IPI) to bring the rest of the cores into SMM before continuing
- We were wrong: The I/O Hub sends the SMI to all cores at once



# Problem summarized

- SMIs make all cores go to SMM at the same time
- TClose is enabled on two logical cores at a time
  - They will read 0xFFs since no device is mapped there
  - Writes to SMM save state will be dropped
- This will make core 1 triple-fault and crash the system



# Tackling the problem

- We had:
  - Control of data fetches on core 0
  - No control of data fetches on core 1
- We tried many things to solve the problem:
  - Finding another device to overlap with the SMM entry point
  - Disabling Simultaneous Multi-Threading (SMT)
  - Sending an INIT IPI / executing SKINIT to ignore SMIs
  - Sending an SMI IPI to trigger an SMI on individual cores

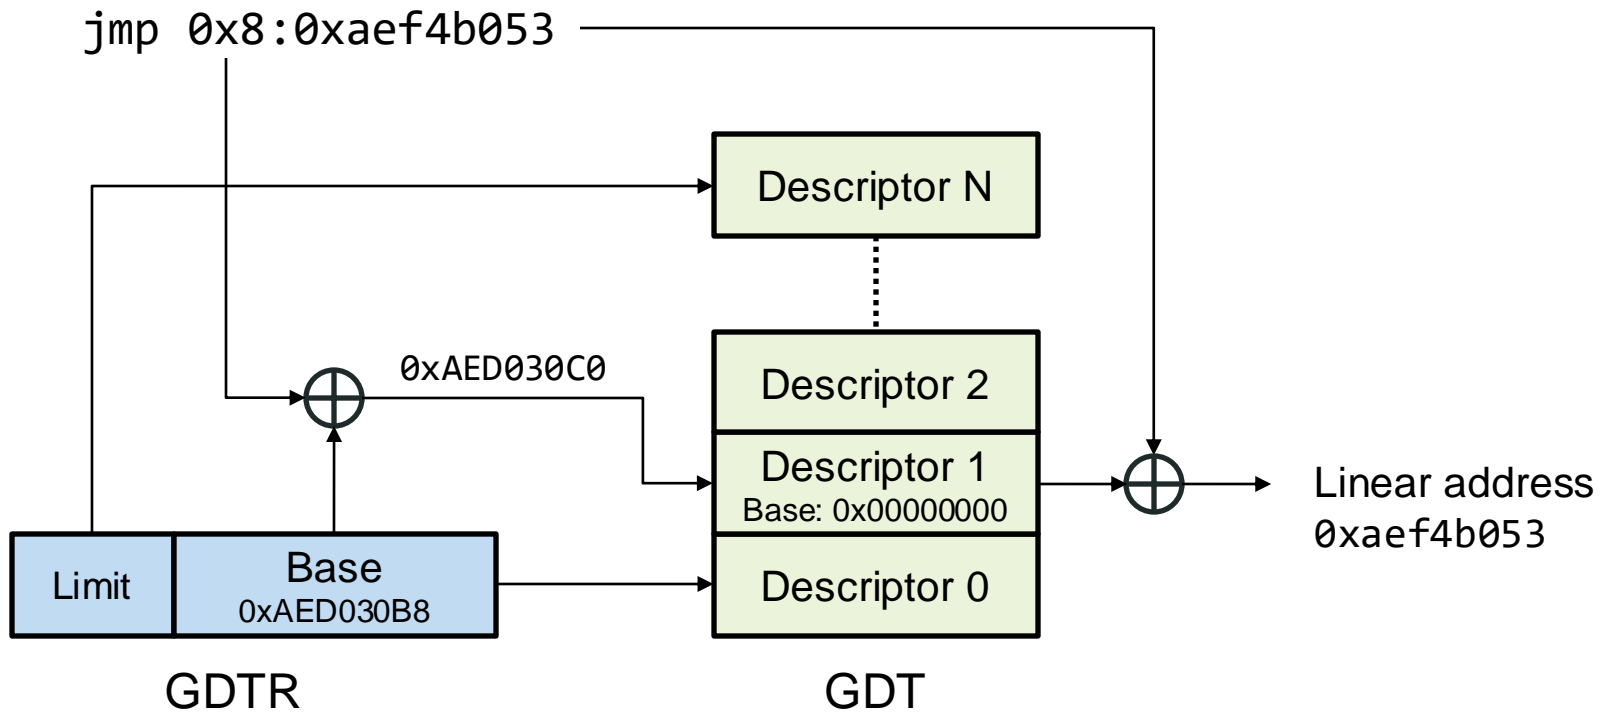


# Running out of options

- Taking a step back:
  - Our lgdt is the issue
  - What happens if the GDTR is loaded all with FFs?
- Let's look into that...



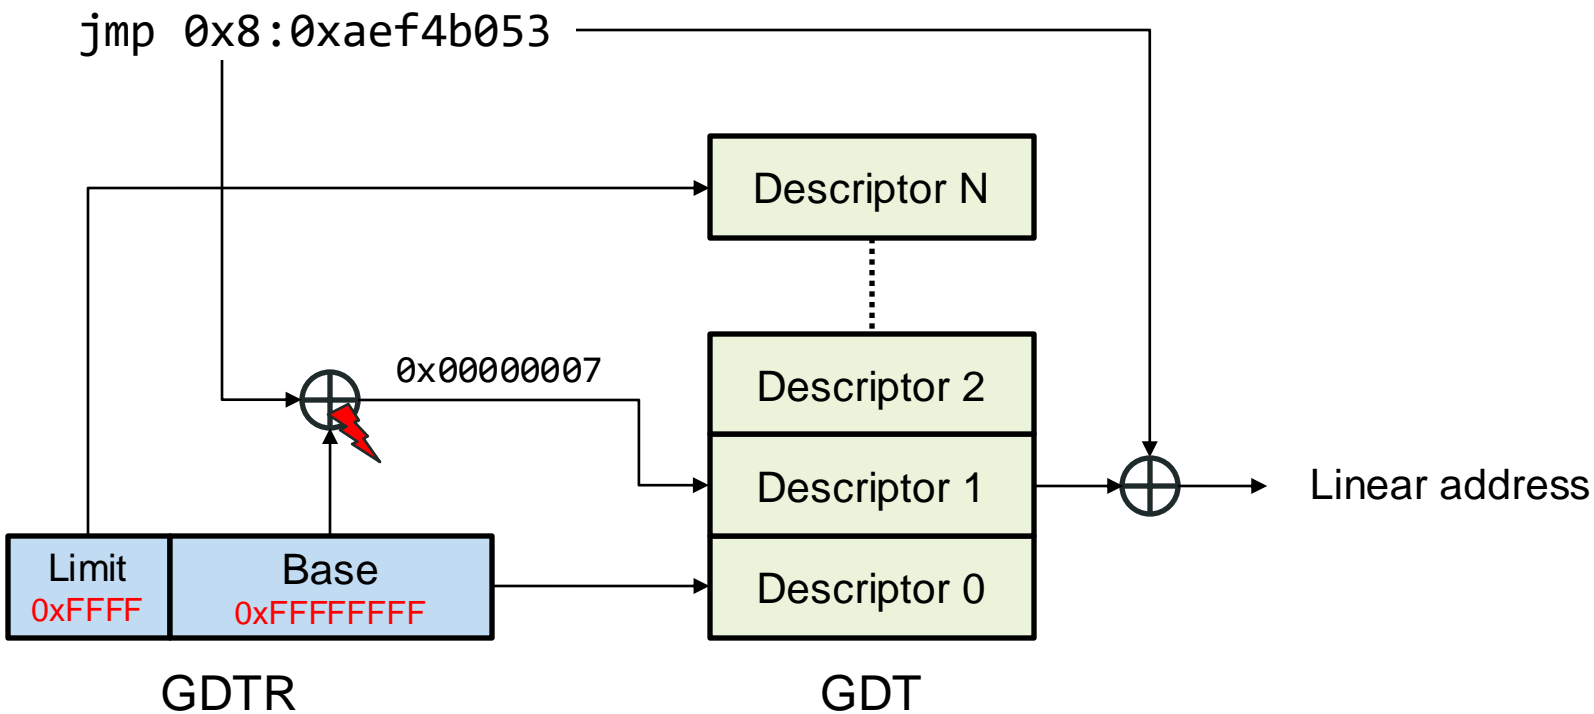
# GDTR wrap-around







# GDTR wrap-around





# Wrap-arounds in x86

- There are two instances of wrap-arounds:
  - The addition between GDT descriptor base and far jmp offset can overflow
  - The addition between the GDTR base and far jmp segment selector can overflow
- We can use the same fake GDT for core 0 and 1
- Added bonus: No need for the SPI BAR remapping




# SMM save state (again)

- For core 0 we use the same technique as before
- For core 1 we:
  - Need to bring core 1 into a known / controlled state
  - We use kernel synchronization APIs to achieve that
    - Deferred Procedure Calls (DPC) on Windows
    - Symmetric Multi-Processing (SMP) on Linux



# Attempt #3




```

mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8 ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053


```



```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```


1. Map + tweak



```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdtd  cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```

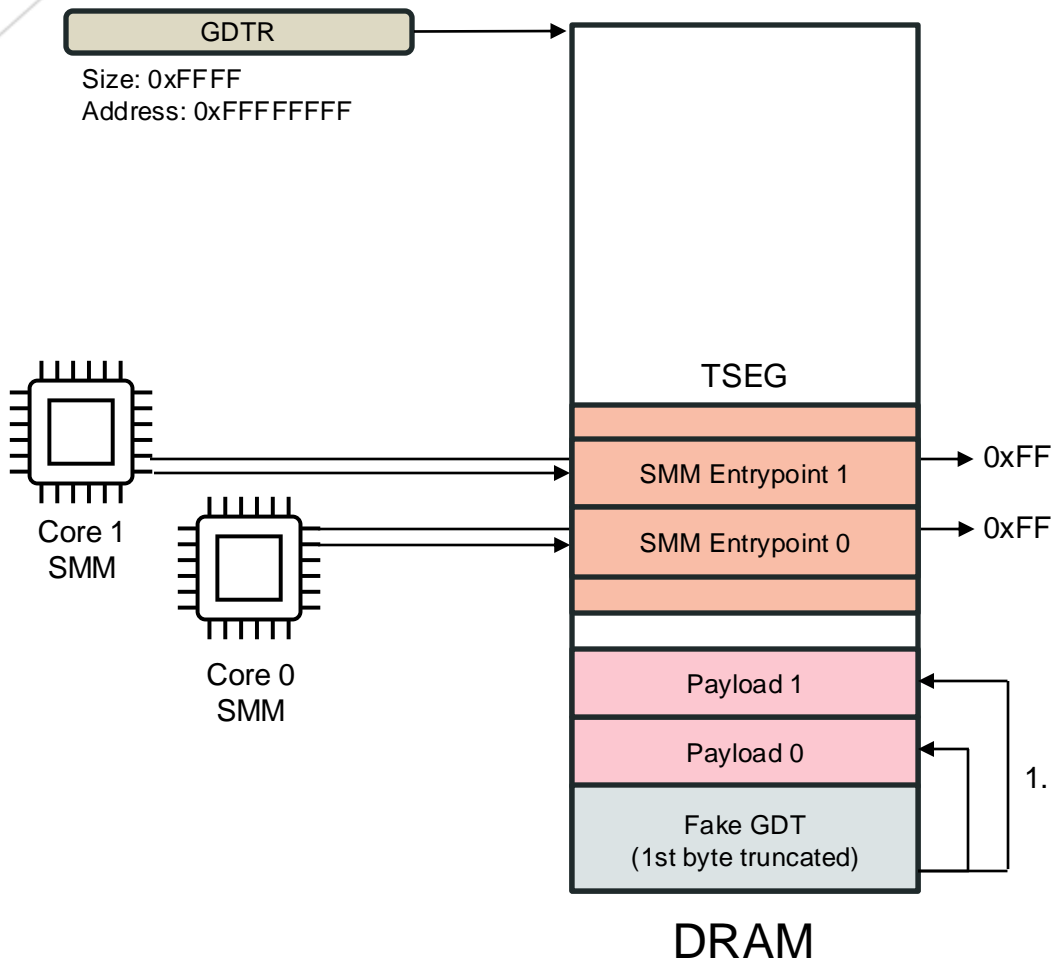
1. Map + tweak



```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdtd  cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

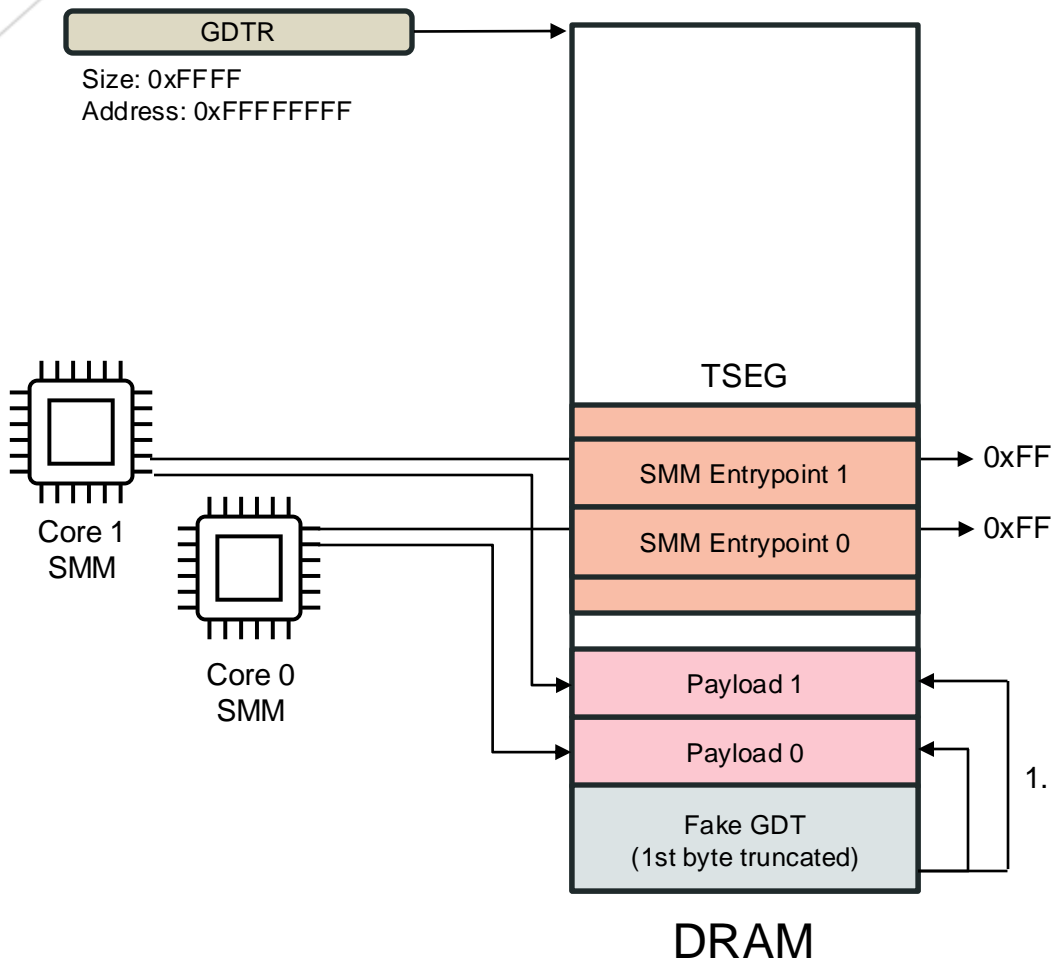
mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```





```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8 ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdt   cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```



```
mov    bx,0x804d    ; 0x8000 + 0x4D
mov    ax,cs:0xfdd8  ; DSC_OFFSET + 0xD8
dec    ax
mov    WORD PTR cs:[bx],ax
mov    eax,cs:0xfdd0 ; DSC_OFFSET + 0xD0
mov    DWORD PTR cs:[bx+0x2],eax
lgdtd  cs:[bx];
mov    ax,0x8
mov    WORD PTR cs:[bx-0x2],ax
mov    edi,0xae43000
lea    eax,[edi+0x8053]

mov    DWORD PTR cs:[bx-0x6],eax
mov    ebx,cr0
and    ebx,0x9ffafff3
or     ebx,0x23
mov    cr0,ebx
jmp    0x8:0xae4b053
```



And it worked!



# Extra steps

- We can execute code in SMM but in protected mode
- Our payload performs the following steps:
  - Reload the GDT to avoid IP misalignments
  - Setup long mode (including page tables)
  - Install an SMI handlers to avoid re-exploiting the issue



# DEMO





# Next attack paths

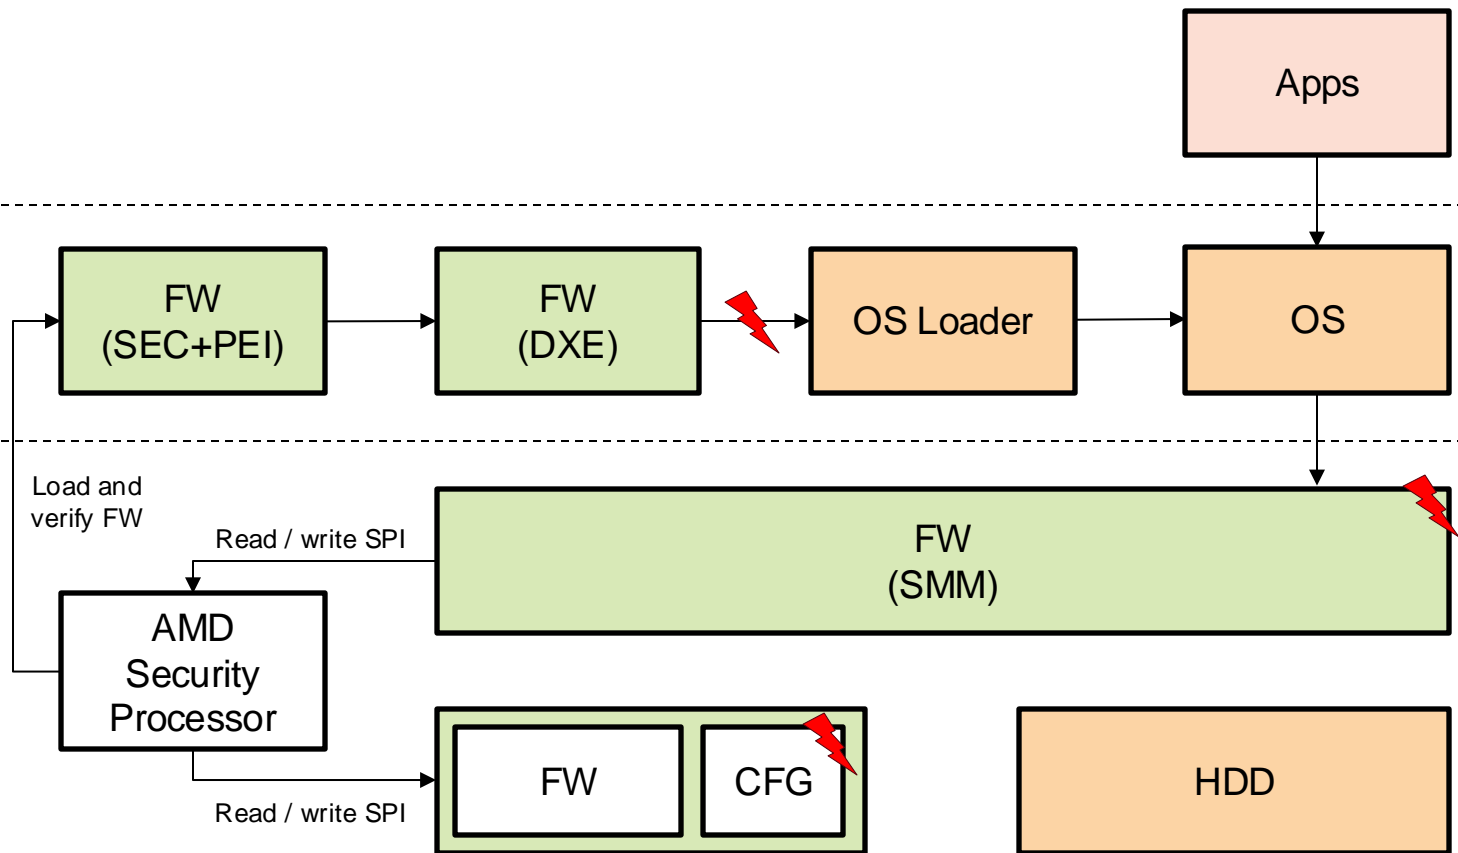
- Next steps depend on the platform configuration
- The firmware is responsible for:
  - Restricting access to the SPI flash (e.g. via ROM Armor)
  - Verifying the firmware chain-of-trust (via Platform Secure Boot)
- If everything is enabled, we can at least break secure boot
- If not, there is potential for firmware implants



Ring 3

Ring 0

Ring -2



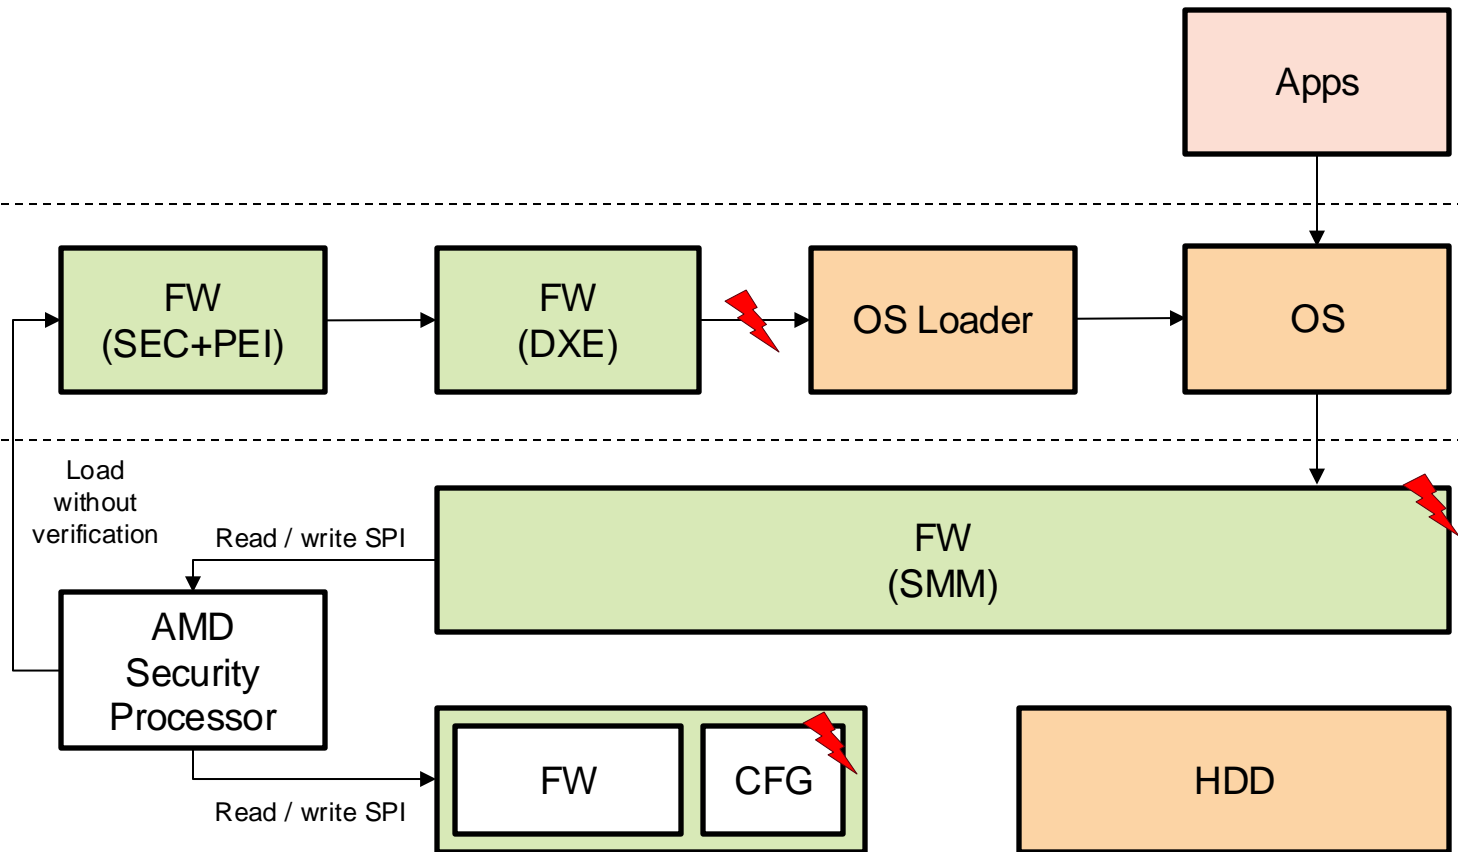




Ring 3

Ring 0

Ring -2

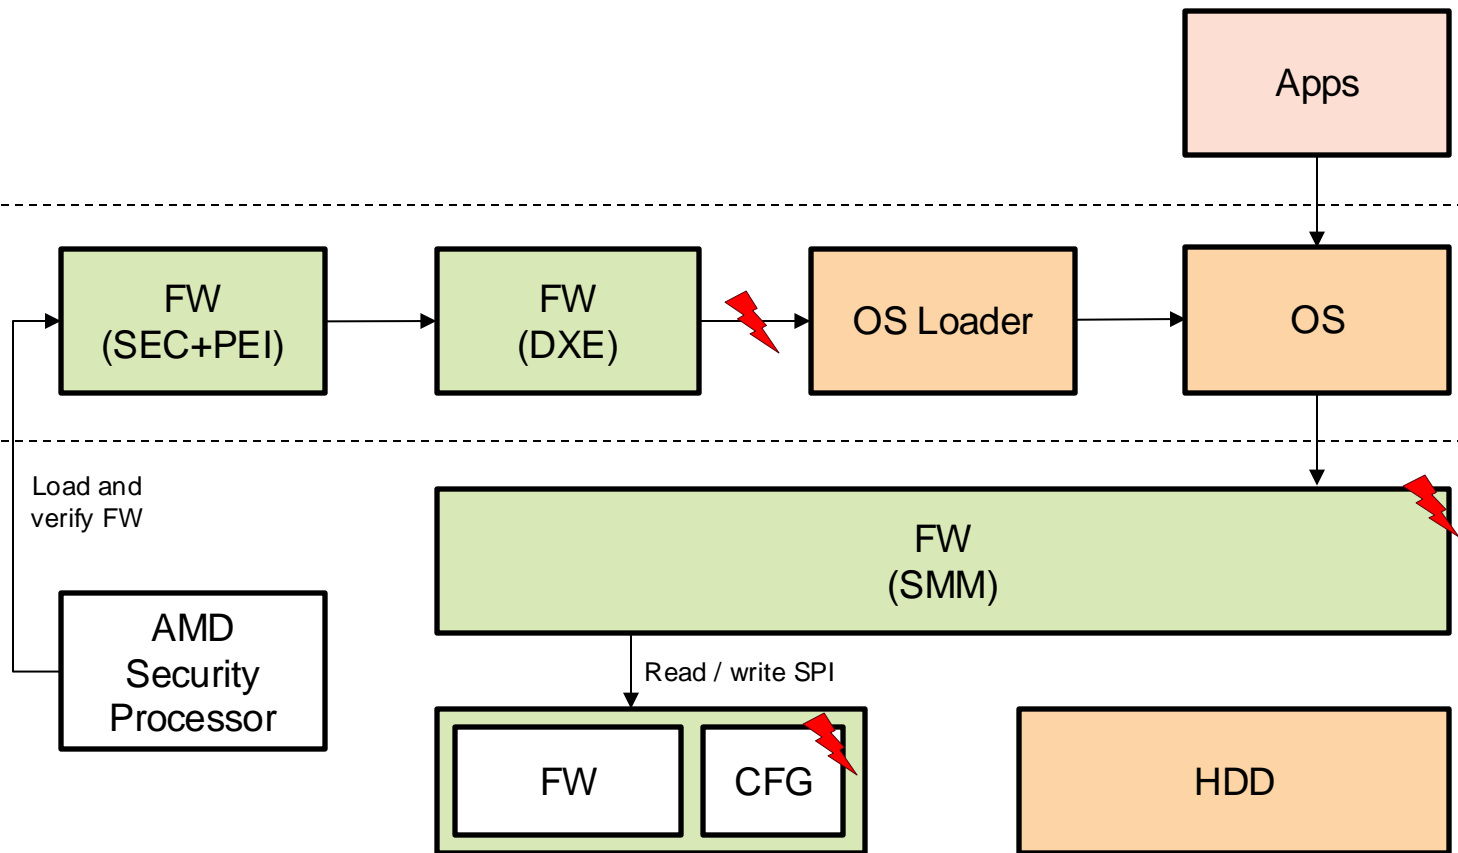




Ring 3

Ring 0

Ring -2

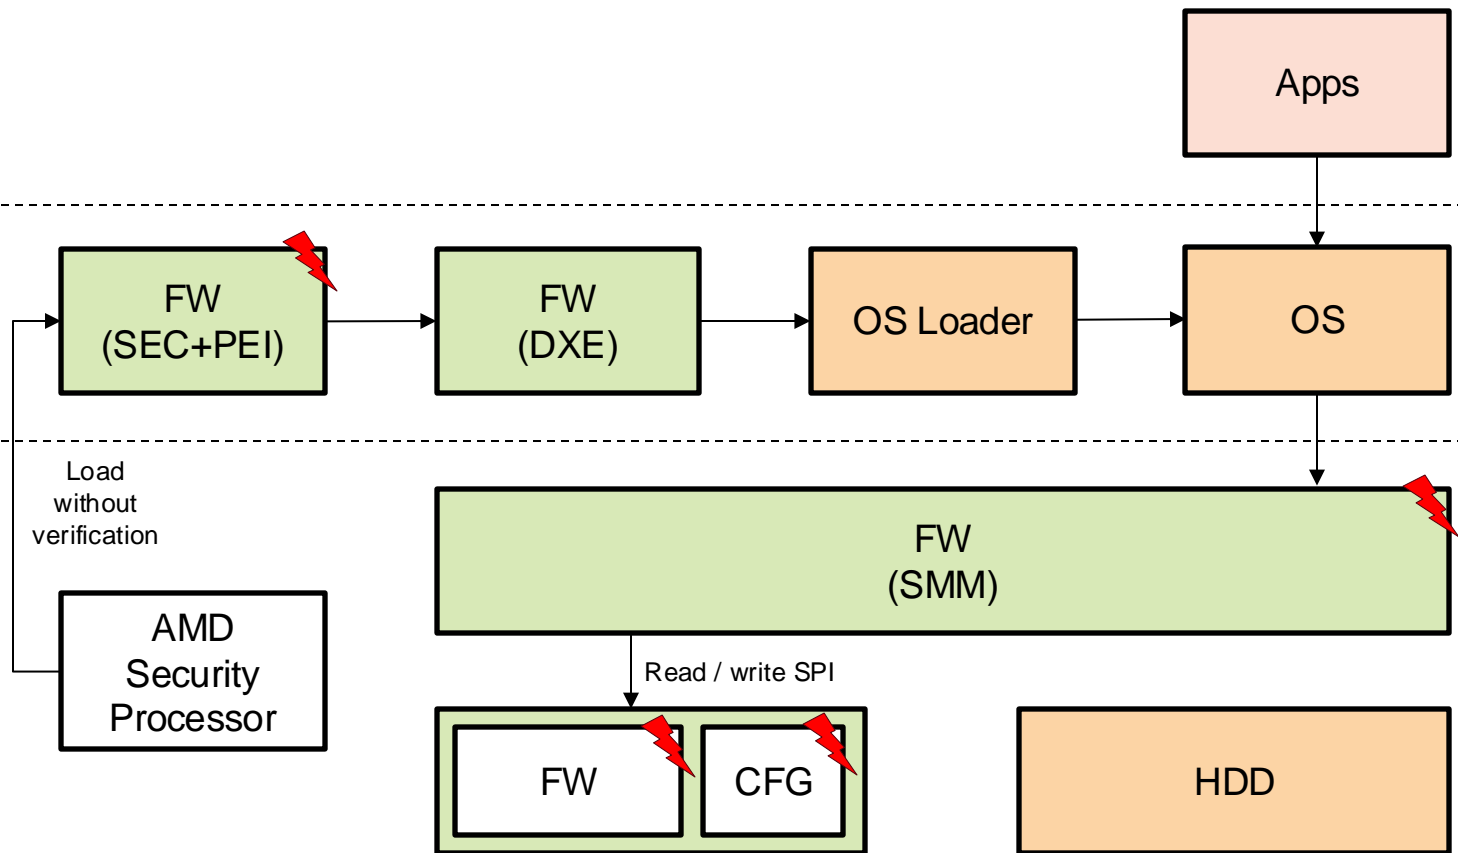




Ring 3

Ring 0

Ring -2





# Platform security (as of 2023)

Vendor	Model	PSB State	ROM Armor State
Acer	Swift 3 SF314-42	Not configured	Not configured
Acer	TravelMate P414-41	Not configured	Configured
ASUS	Strix G513QR	Not configured	Not configured
Lenovo	Thinkpad P16s	Configured*	Not configured
Lenovo	IdeaPad 1	Not configured	Not configured
Lenovo	Thinkpad T495s	Not configured	Not configured
Huawei	Matebook D16	Not configured	Not configured
HP	15s	Not configured	Not configured
Microsoft	Surface 4	Configured	Unknown
MSI	Bravo 15	Not configured	Not configured



# Platform security continued

- In previous research we discovered that the PSB can be permanently disabled by burning specific fuses:

PSB Status	PSB_EN	CUSTOMER_KEY_LOCK
Not Enabled	0	0
Enabled	1	1
Disabled	0	1

- Once a system is compromised, doing this leaves it vulnerable to firmware implants forever



# Outro



# Affected systems

- Pretty much all of them
  - Ryzen series
  - Ryzen Threadripper series
  - EPYC series
- Total number of affected chips: 100s of millions
- AMD advisory AMD-SB-7014 published at <https://www.amd.com/en/resources/product-security/bulletin/amd-sb-7014.html>



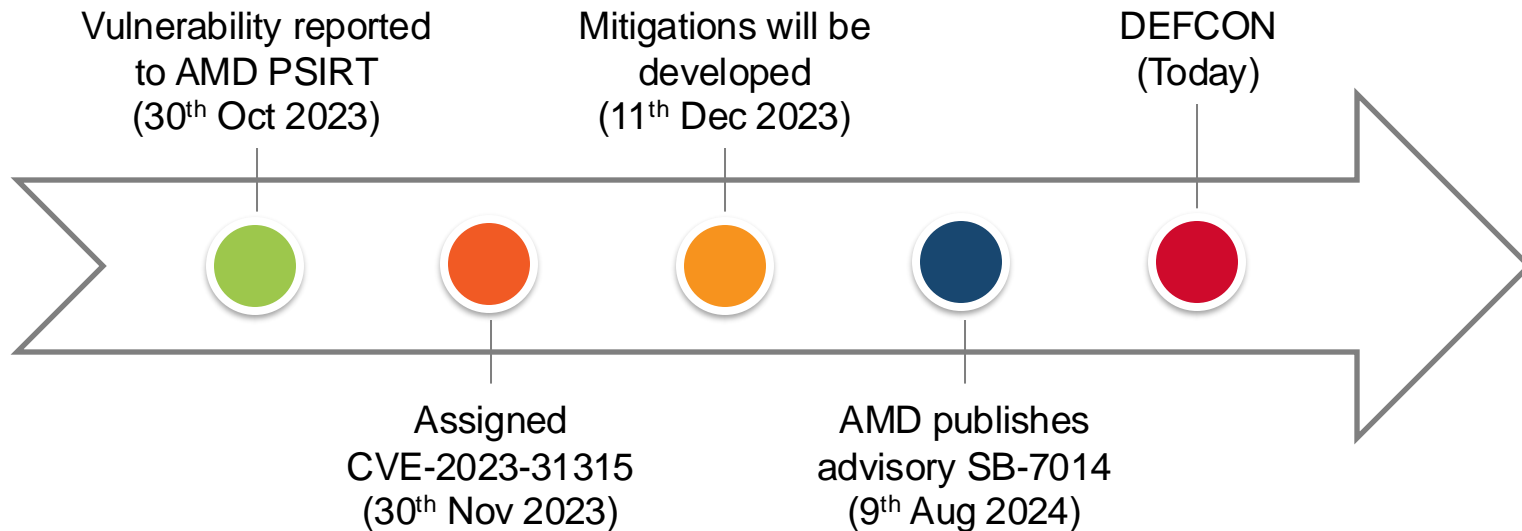
# Mitigations

- AMD:
  - A microcode update is available
  - Con: Might not cover all affected systems due to product EOL
- OEMs:
  - Modify SMM entry point code to detect if TClose bit is enabled and abort execution
  - Can be done at the reference code level
  - Con: Specific to one OEM or even specific systems
- Users:
  - A hypervisor could be used to trap accesses on the TSEG mask MSR





# Timeline





# AMD's response

*AMD thanks IOActive for identifying the vulnerability and working with AMD to protect end-users.*

*AMD has identified and deployed mitigations for this vulnerability. A full list of impacted products and mitigation options is available in our product security bulletin AMD-SB-7014 which may be found here:*

*<https://www.amd.com/productsecurity>*

*AMD welcomes collaboration with the security community and encourages researchers to submit their findings to AMD PSIRT using the product security page above.*



# Conclusions

- The vulnerability has been around for nearly two decades
- The complexity of modern architectures plays in favor of attackers
- The flexibility of segmentation played a crucial role for exploitation
- Exploitation requires in-depth understanding of the architecture
- This issue can be exploited without requiring physical presence

*Exploit code will be released soon*

*Stay tuned!*



# Questions?